

WIKIPEDIA

Artificial intelligence

Artificial intelligence (AI) is intelligence demonstrated by machines, as opposed to the **natural intelligence** displayed by animals and humans. AI research has been defined as the field of study of intelligent agents, which refers to any system that perceives its environment and takes actions that maximize its chance of achieving its goals.^[a]

The term "artificial intelligence" had previously been used to describe machines that mimic and display "human" cognitive skills that are associated with the human mind, such as "learning" and "problem-solving". This definition has since been rejected by major AI researchers who now describe AI in terms of rationality and acting rationally, which does not limit how intelligence can be articulated.^[b]

AI applications include advanced web search engines (e.g., Google), recommendation systems (used by YouTube, Amazon and Netflix), understanding human speech (such as Siri and Alexa), self-driving cars (e.g., Tesla), automated decision-making and competing at the highest level in strategic game systems (such as chess and Go).^[2] As machines become increasingly capable, tasks considered to require "intelligence" are often removed from the definition of AI, a phenomenon known as the AI effect.^[3] For instance, optical character recognition is frequently excluded from things considered to be AI,^[4] having become a routine technology.^[5]

Artificial intelligence was founded as an academic discipline in 1956, and in the years since has experienced several waves of optimism,^{[6][7]} followed by disappointment and the loss of funding (known as an "AI winter"),^{[8][9]} followed by new approaches, success and renewed funding.^{[7][10]} AI research has tried and discarded many different approaches since its founding, including simulating the brain, modeling human problem solving, formal logic, large databases of knowledge and imitating animal behavior. In the first decades of the 21st century, highly mathematical-statistical machine learning has dominated the field, and this technique has proved highly successful, helping to solve many challenging problems throughout industry and academia.^{[10][11]}

The various sub-fields of AI research are centered around particular goals and the use of particular tools. The traditional goals of AI research include reasoning, knowledge representation, planning, learning, natural language processing, perception, and the ability to move and manipulate objects.^[c] General intelligence (the ability to solve an arbitrary problem) is among the field's long-term goals.^[12] To solve these problems, AI researchers have adapted and integrated a wide range of problem-solving techniques – including search and mathematical optimization, formal logic, artificial neural networks, and methods based on statistics, probability and economics. AI also draws upon computer science, psychology, linguistics, philosophy, and many other fields.

The field was founded on the assumption that human intelligence "can be so precisely described that a machine can be made to simulate it".^[d] This raised philosophical arguments about the mind and the ethical consequences of creating artificial beings endowed with human-like intelligence; these issues have previously been explored by myth, fiction and philosophy since antiquity.^[14] Computer scientists and philosophers have since suggested that AI may become an existential risk to humanity if its rational capacities are not steered towards beneficial goals.^[e]

Contents



History

Goals

[Reasoning, problem-solving](#)
[Knowledge representation](#)
[Planning](#)
[Learning](#)
[Natural language processing](#)
[Perception](#)
[Motion and manipulation](#)
[Social intelligence](#)
[General intelligence](#)

Tools

[Search and optimization](#)
[Logic](#)
[Probabilistic methods for uncertain reasoning](#)
[Classifiers and statistical learning methods](#)
[Artificial neural networks](#)
[Deep learning](#)
[Specialized languages and hardware](#)

Applications

Philosophy

[Defining artificial intelligence](#)
[Evaluating approaches to AI](#)
[Symbolic AI and its limits](#)
[Neat vs. scruffy](#)
[Soft vs. hard computing](#)
[Narrow vs. general AI](#)
[Machine consciousness, sentience and mind](#)
[Consciousness](#)
[Computationalism and functionalism](#)
[Robot rights](#)

Future

[Superintelligence](#)
[Risks](#)
[Technological unemployment](#)
[Bad actors and weaponized AI](#)
[Algorithmic bias](#)
[Existential risk](#)
[Other issues: Copyright](#)
[Ethical machines](#)
[Regulation](#)

In fiction

Scientific diplomacy

[Warfare](#)
[Russo-Ukrainian War](#)
[Warfare regulations](#)

Cybersecurity

[Czech Republic's approach](#)

[Germany's approach](#)

[European Union's approach](#)

[Russo-Ukrainian War](#)

Election security

Future of work

[Facial recognition](#)

[AI and school](#)

[AI and medicine](#)

[AI in business](#)

[Business and diplomacy](#)

AI and foreign policy

See also

Explanatory notes

Citations

References

[AI textbooks](#)

[History of AI](#)

[Other sources](#)

Further reading

External links

History

Artificial beings with intelligence appeared as [storytelling devices](#) in antiquity,^[15] and have been common in fiction, as in [Mary Shelley's *Frankenstein*](#) or [Karel Čapek's *R.U.R.*](#)^[16] These characters and their fates raised many of the same issues now discussed in the [ethics of artificial intelligence](#).^[17]

The study of mechanical or "formal" reasoning began with [philosophers](#) and [mathematicians](#) in antiquity. The study of mathematical logic led directly to [Alan Turing's theory of computation](#), which suggested that a machine, by shuffling symbols as simple as "0" and "1", could simulate any conceivable act of mathematical deduction. This insight that digital computers can simulate any process of formal reasoning is known as the [Church–Turing thesis](#).^[18] This, along with concurrent discoveries in [neurobiology](#), [information theory](#) and [cybernetics](#), led researchers to consider the possibility of building an electronic brain.^[19] The first work that is now generally recognized as AI was [McCullouch and Pitts' 1943 formal design for Turing-complete "artificial neurons"](#).^[20]



Silver [didrachma](#) from [Crete](#) depicting [Talos](#), an ancient mythical [automaton](#) with [artificial intelligence](#)

By the 1950s, two visions for how to achieve machine intelligence emerged. One vision, known as [Symbolic AI](#) or [GOFAl](#), was to use computers to create a symbolic representation of the world and systems that could reason about the world. Proponents included [Allen Newell](#), [Herbert A. Simon](#), and [Marvin Minsky](#). Closely associated with this approach was the ["heuristic search"](#) approach,

which likened intelligence to a probability space. The second vision, known as the connectionist approach, sought to achieve intelligence through learning. Proponents of this approach, most prominently Frank Rosenblatt, sought to connect Perceptron in ways inspired by connections of neurons.^[21] James Manyika and others have compared the two approaches to the mind (Symbolic AI) and the brain (connectionist). Manyika argues that symbolic approaches dominated the push for artificial intelligence in this period, due in part to its connection to intellectual traditions of Descartes, Boole, Gottlob Frege, Bertrand Russell, and others. Connectionist approaches based on cybernetics or artificial neural networks were pushed to the background but have gained new prominence in recent decades.^[22]

The field of AI research was born at a workshop at Dartmouth College in 1956.^{[f][25]} The attendees became the founders and leaders of AI research.^[g] They and their students produced programs that the press described as "astonishing":^[h] computers were learning checkers strategies, solving word problems in algebra, proving logical theorems and speaking English.^{[i][27]} By the middle of the 1960s, research in the U.S. was heavily funded by the Department of Defense^[28] and laboratories had been established around the world.^[29]

Researchers in the 1960s and the 1970s were convinced that symbolic approaches would eventually succeed in creating a machine with artificial general intelligence and considered this the goal of their field.^[30] Herbert Simon predicted, "machines will be capable, within twenty years, of doing any work a man can do".^[31] Marvin Minsky agreed, writing, "within a generation ... the problem of creating 'artificial intelligence' will substantially be solved".^[32] They had failed to recognize the difficulty of some of the remaining tasks. Progress slowed and in 1974, in response to the criticism of Sir James Lighthill^[33] and ongoing pressure from the US Congress to fund more productive projects, both the U.S. and British governments cut off exploratory research in AI. The next few years would later be called an "AI winter", a period when obtaining funding for AI projects was difficult.^[8]

In the early 1980s, AI research was revived by the commercial success of expert systems,^[34] a form of AI program that simulated the knowledge and analytical skills of human experts. By 1985, the market for AI had reached over a billion dollars. At the same time, Japan's fifth generation computer project inspired the U.S. and British governments to restore funding for academic research.^[7] However, beginning with the collapse of the Lisp Machine market in 1987, AI once again fell into disrepute, and a second, longer-lasting winter began.^[9]

Many researchers began to doubt that the symbolic approach would be able to imitate all the processes of human cognition, especially perception, robotics, learning and pattern recognition. A number of researchers began to look into "sub-symbolic" approaches to specific AI problems.^[35] Robotics researchers, such as Rodney Brooks, rejected symbolic AI and focused on the basic engineering problems that would allow robots to move, survive, and learn their environment.^[i] Interest in neural networks and "connectionism" was revived by Geoffrey Hinton, David Rumelhart and others in the middle of the 1980s.^[40] Soft computing tools were developed in the 1980s, such as neural networks, fuzzy systems, Grey system theory, evolutionary computation and many tools drawn from statistics or mathematical optimization.

AI gradually restored its reputation in the late 1990s and early 21st century by finding specific solutions to specific problems. The narrow focus allowed researchers to produce verifiable results, exploit more mathematical methods, and collaborate with other fields (such as statistics, economics and mathematics).^[41] By 2000, solutions developed by AI researchers were being widely used, although in the 1990s they were rarely described as "artificial intelligence".^[11]

Faster computers, algorithmic improvements, and access to large amounts of data enabled advances in machine learning and perception; data-hungry deep learning methods started to dominate accuracy benchmarks around 2012.^[42] According to Bloomberg's Jack Clark, 2015 was a

landmark year for artificial intelligence. Google increased from a "sporadic user" to one of the largest users of AI. This has led to an increase in affordable neural networks, due to a rise in cloud computing infrastructure and to an increase in research tools and datasets.^[10] In a 2017 survey, one in five companies reported they had "incorporated AI in some offerings or processes".^[43] The amount of research into AI (measured by total publications) increased by 50% in the years 2015–2019.^[44]

Numerous academic researchers became concerned that AI was no longer pursuing the original goal of creating versatile, fully intelligent machines. Much of current research involves statistical AI, which is overwhelmingly used to solve specific problems, even highly successful techniques such as deep learning. This concern has led to the subfield of artificial general intelligence (or "AGI"), which had several well-funded institutions by the 2010s.^[12]

Goals

The general problem of simulating (or creating) intelligence has been broken down into subproblems. These consist of particular traits or capabilities that researchers expect an intelligent system to display. The traits described below have received the most attention.^[c]

Reasoning, problem-solving

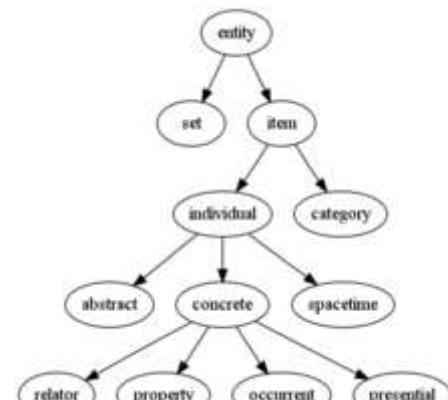
Early researchers developed algorithms that imitated step-by-step reasoning that humans use when they solve puzzles or make logical deductions.^[45] By the late 1980s and 1990s, AI research had developed methods for dealing with uncertain or incomplete information, employing concepts from probability and economics.^[46]

Many of these algorithms proved to be insufficient for solving large reasoning problems because they experienced a "combinatorial explosion": they became exponentially slower as the problems grew larger.^[47] Even humans rarely use the step-by-step deduction that early AI research could model. They solve most of their problems using fast, intuitive judgments.^[48]

Knowledge representation

Knowledge representation and knowledge engineering^[49] allow AI programs to answer questions intelligently and make deductions about real-world facts.

A representation of "what exists" is an ontology: the set of objects, relations, concepts, and properties formally described so that software agents can interpret them.^[50] The most general ontologies are called upper ontologies, which attempt to provide a foundation for all other knowledge and act as mediators between domain ontologies that cover specific knowledge about a particular knowledge domain (field of interest or area of concern). A truly intelligent program would also need access to commonsense knowledge; the set of facts that an average person knows. The semantics of an ontology is typically represented in description logic, such as the Web Ontology Language.^[51]



An ontology represents knowledge as a set of concepts within a domain and the relationships between those concepts.

AI research has developed tools to categories and relations between objects, such as objects, properties, states and time; causes and effects; knowledge about knowledge (what we know about what other people know); default reasoning (things that humans assume are true until they are told differently and will remain true even when other facts are changing); as well as other domains. Among the most difficult problems in AI are: the breadth of commonsense knowledge (the number of atomic facts that the average person knows is enormous); and the sub-symbolic form of most commonsense knowledge (much of what people know is not represented as "facts" or "statements" that they could express verbally).^[48]

Formal knowledge representations are used in content-based indexing and retrieval,^[57] scene interpretation,^[58] clinical decision support,^[59] knowledge discovery (mining "interesting" and actionable inferences from large databases),^[60] and other areas.^[61]

Planning

An intelligent agent that can plan makes a representation of the state of the world, makes predictions about how their actions will change it and make choices that maximize the utility (or "value") of the available choices.^[62] In classical planning problems, the agent can assume that it is the only system acting in the world, allowing the agent to be certain of the consequences of its actions.^[63] However, if the agent is not the only actor, then it requires that the agent reason under uncertainty, and continuously re-assess its environment and adapt.^[64] Multi-agent planning uses the cooperation and competition of many agents to achieve a given goal. Emergent behavior such as this is used by evolutionary algorithms and swarm intelligence.^[65]

Learning

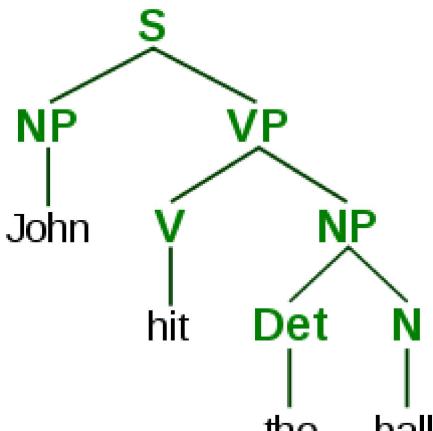
Machine learning (ML), a fundamental concept of AI research since the field's inception,^[1] is the study of computer algorithms that improve automatically through experience.^[m]

Unsupervised learning finds patterns in a stream of input. Supervised learning requires a human to label the input data first, and comes in two main varieties: classification and numerical regression. Classification is used to determine what category something belongs in – the program sees a number of examples of things from several categories and will learn to classify new inputs. Regression is the attempt to produce a function that describes the relationship between inputs and outputs and predicts how the outputs should change as the inputs change. Both classifiers and regression learners can be viewed as "function approximators" trying to learn an unknown (possibly implicit) function; for example, a spam classifier can be viewed as learning a function that maps from the text of an email to one of two categories, "spam" or "not spam".^[69] In reinforcement learning the agent is rewarded for good responses and punished for bad ones. The agent classifies its responses to form a strategy for operating in its problem space.^[70] Transfer learning is when the knowledge gained from one problem is applied to a new problem.^[71]

Computational learning theory can assess learners by computational complexity, by sample complexity (how much data is required), or by other notions of optimization.^[72]

Natural language processing

Natural language processing (NLP)^[73] allows machines to read and understand human language. A sufficiently powerful natural language processing system would enable natural-language user interfaces and the acquisition of knowledge directly from human-written sources, such as



A parse tree represents the syntactic structure of a sentence according to some formal grammar.

ne in straightforward applications of NLP val, question answering and machine translation.^[44]

Symbolic AI used formal syntax to translate the deep structure of sentences into logic. This failed to produce useful applications, due to the intractability of logic^[47] and the breadth of commonsense knowledge.^[56] Modern statistical techniques include co-occurrence frequencies (how often one word appears near another), "Keyword spotting" (searching for a particular word to retrieve information), transformer-based deep learning (which finds patterns in text), and others.^[75] They have achieved acceptable accuracy at the page or paragraph level, and, by 2019, could generate coherent text.^[76]

Perception

Machine perception^[77] is the ability to use input from sensors (such as cameras, microphones, wireless signals, and active lidar, sonar, radar, and tactile sensors) to deduce aspects of the world. Applications include speech recognition,^[78] facial recognition, and object recognition.^[79] Computer vision is the ability to analyze visual input.^[80]



Feature detection (pictured: edge detection) helps AI compose informative abstract structures out of raw data.

Motion and manipulation

AI is heavily used in robotics.^[81] Localization is how a robot knows its location and maps its environment. When given a small, static, and visible environment, this is easy; however, dynamic environments, such as (in endoscopy) the interior of a patient's breathing body, pose a greater challenge.^[82]

Motion planning is the process of breaking down a movement task into "primitives" such as individual joint movements. Such movement often involves compliant motion, a process where movement requires maintaining physical contact with an object. Robots can learn from experience how to move efficiently despite the presence of friction and gear slippage.^[83]

Social intelligence

Affective computing is an interdisciplinary umbrella that comprises systems that recognize, interpret, process or simulate human feeling, emotion and mood.^[85] For example, some virtual assistants are programmed to speak conversationally or even to banter humorously; it makes them appear more sensitive to the emotional dynamics of human interaction, or to otherwise facilitate human-computer interaction. However, this tends to give naïve users an unrealistic conception of how intelligent existing computer agents actually are.^[86] Moderate successes related to affective computing include textual sentiment analysis and, more recently, multimodal sentiment analysis, wherein AI classifies the affects displayed by a videotaped subject.^[87]

General intelligence



A machine with general intelligence can solve a wide variety of problems with breadth and versatility similar to human intelligence. There are several competing ideas about how to develop artificial general intelligence. [Hans Moravec](#) and [Marvin Minsky](#) argue that work in different individual domains can be incorporated into an advanced [multi-agent system](#) or cognitive architecture with general intelligence.^[88] [Pedro Domingos](#) hopes that there is a conceptually straightforward, but mathematically difficult, "master algorithm" that could lead to AGI.^[89] Others believe that [anthropomorphic](#) features like an [artificial brain](#)^[90] or simulated [child development](#)^[n] will someday reach a critical point where general intelligence emerges.



Kismet, a robot with rudimentary social skills^[84]

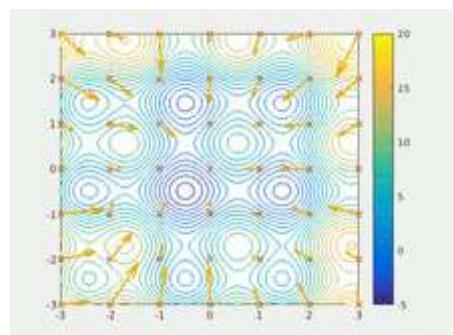
Tools

Search and optimization

Many problems in AI can be solved theoretically by intelligently searching through many possible solutions:^[91] Reasoning can be reduced to performing a search. For example, logical proof can be viewed as searching for a path that leads from [premises](#) to [conclusions](#), where each step is the application of an [inference rule](#).^[92] Planning algorithms search through trees of goals and subgoals, attempting to find a path to a target goal, a process called [means-ends analysis](#).^[93] Robotics algorithms for moving limbs and grasping objects use [local searches](#) in [configuration space](#).^[94]

Simple exhaustive searches^[95] are rarely sufficient for most real-world problems: the [search space](#) (the number of places to search) quickly grows to [astronomical numbers](#). The result is a search that is [too slow](#) or never completes. The solution, for many problems, is to use "[heuristics](#)" or "[rules of thumb](#)" that prioritize choices in favor of those more likely to reach a goal and to do so in a shorter number of steps. In some search methodologies, heuristics can also serve to eliminate some choices unlikely to lead to a goal (called "[pruning the search tree](#)"). Heuristics supply the program with a "[best guess](#)" for the path on which the solution lies.^[96] Heuristics limit the search for solutions into a smaller sample size.^[97]

A very different kind of search came to prominence in the 1990s, based on the mathematical theory of [optimization](#). For many problems, it is possible to begin the search with some form of a guess and then refine the guess incrementally until no more refinements can be made. These algorithms can be visualized as blind [hill climbing](#): we begin the search at a random point on the landscape, and then, by jumps or steps, we keep moving our guess uphill, until we reach the top. Other related optimization algorithms include [random optimization](#), [beam search](#) and [metaheuristics](#) like [simulated annealing](#).^[98] Evolutionary computation uses a form of optimization search. For example, they may begin with a population of organisms (the guesses) and then allow them to mutate and recombine, selecting only the fittest to survive each generation (refining the guesses). Classic [evolutionary algorithms](#) include [genetic algorithms](#), [gene expression programming](#), and [genetic](#)



A particle swarm seeking the global minimum

programming.^[99] Alternatively, different processes can coordinate via swarm intelligence algorithms. Two popular optimization (inspired by bird flocking) and ant colony optimization (inspired by ant trails).^[100]

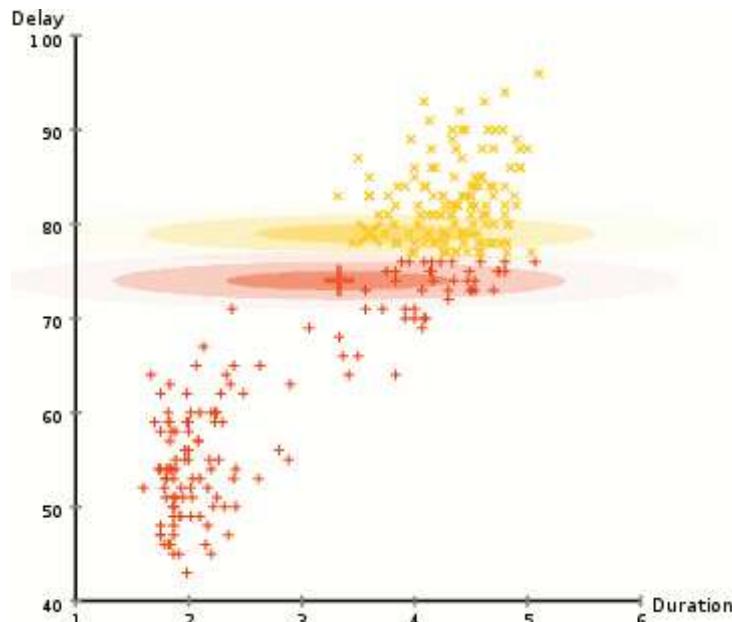
Logic

Logic^[101] is used for knowledge representation and problem-solving, but it can be applied to other problems as well. For example, the satplan algorithm uses logic for planning^[102] and inductive logic programming is a method for learning.^[103]

Several different forms of logic are used in AI research. Propositional logic^[104] involves truth functions such as "or" and "not". First-order logic^[105] adds quantifiers and predicates and can express facts about objects, their properties, and their relations with each other. Fuzzy logic assigns a "degree of truth" (between 0 and 1) to vague statements such as "Alice is old" (or rich, or tall, or hungry), that are too linguistically imprecise to be completely true or false.^[106] Default logics, non-monotonic logics and circumscription are forms of logic designed to help with default reasoning and the qualification problem.^[55] Several extensions of logic have been designed to handle specific domains of knowledge, such as description logics,^[51] situation calculus, event calculus and fluent calculus (for representing events and time);^[52] causal calculus,^[53] belief calculus (belief revision); and modal logics.^[54] Logics to model contradictory or inconsistent statements arising in multi-agent systems have also been designed, such as paraconsistent logics.^[107]

Probabilistic methods for uncertain reasoning

Many problems in AI (including in reasoning, planning, learning, perception, and robotics) require the agent to operate with incomplete or uncertain information. AI researchers have devised a number of tools to solve these problems using methods from probability theory and economics.^[108] Bayesian networks^[109] are a very general tool that can be used for various problems, including reasoning (using the Bayesian inference algorithm),^[110] learning (using the expectation-maximization algorithm),^[111] planning (using decision networks),^[112] and perception (using dynamic Bayesian networks).^[113] Probabilistic algorithms can also be used for filtering, prediction, smoothing and finding explanations for streams of data, helping perception systems to analyze processes that occur over time (e.g., hidden Markov models or Kalman filters).^[115]



Expectation-maximization clustering of Old Faithful eruption data starts from a random guess but then successfully converges on an accurate clustering of the two physically distinct modes of eruption.

A key concept from the science of economics is "utility", a measure of how valuable something is to an intelligent agent. Precise mathematical tools have been developed that analyze how an agent can make choices and plan, using decision theory, decision analysis,^[116] and information value

theory.^[117] These tools include mode' networks,^[115] game theory and mech.

sion processes,^[118] dynamic decision



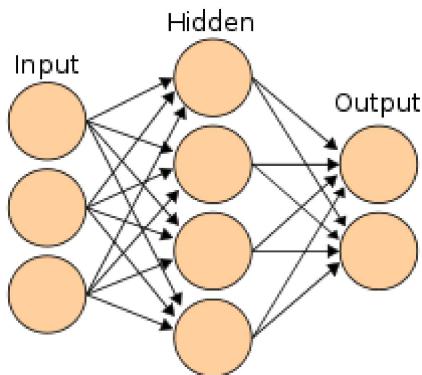
Classifiers and statistical learning methods

The simplest AI applications can be divided into two types: classifiers ("if shiny then diamond") and controllers ("if diamond then pick up"). Controllers do, however, also classify conditions before inferring actions, and therefore classification forms a central part of many AI systems. Classifiers are functions that use pattern matching to determine the closest match. They can be tuned according to examples, making them very attractive for use in AI. These examples are known as observations or patterns. In supervised learning, each pattern belongs to a certain predefined class. A class is a decision that has to be made. All the observations combined with their class labels are known as a data set. When a new observation is received, that observation is classified based on previous experience.^[120]

A classifier can be trained in various ways; there are many statistical and machine learning approaches. The decision tree is the simplest and most widely used symbolic machine learning algorithm.^[121] K-nearest neighbor algorithm was the most widely used analogical AI until the mid-1990s.^[122] Kernel methods such as the support vector machine (SVM) displaced k-nearest neighbor in the 1990s.^[123] The naive Bayes classifier is reportedly the "most widely used learner"^[124] at Google, due in part to its scalability.^[125] Neural networks are also used for classification.^[126]

Classifier performance depends greatly on the characteristics of the data to be classified, such as the dataset size, distribution of samples across classes, dimensionality, and the level of noise. Model-based classifiers perform well if the assumed model is an extremely good fit for the actual data. Otherwise, if no matching model is available, and if accuracy (rather than speed or scalability) is the sole concern, conventional wisdom is that discriminative classifiers (especially SVM) tend to be more accurate than model-based classifiers such as "naive Bayes" on most practical data sets.^[127]

Artificial neural networks



A neural network is an interconnected group of nodes, akin to the vast network of neurons in the human brain.

Neural networks^[126] were inspired by the architecture of neurons in the human brain. A simple "neuron" N accepts input from other neurons, each of which, when activated (or "fired"), casts a weighted "vote" for or against whether neuron N should itself activate. Learning requires an algorithm to adjust these weights based on the training data; one simple algorithm (dubbed "fire together, wire together") is to increase the weight between two connected neurons when the activation of one triggers the successful activation of another. Neurons have a continuous spectrum of activation; in addition, neurons can process inputs in a nonlinear way rather than weighing straightforward votes.

Modern neural networks model complex relationships between inputs and outputs and find patterns in data. They can learn continuous functions and even digital logical operations.

Neural networks can be viewed as a type of mathematical optimization – they perform gradient descent on a multi-dimensional topology that was created by training the network. The most

common training technique is the backpropagation algorithm.^[128] Other learning techniques for neural networks are Hebbian learning ("learning from experience together"), GMDH or competitive learning.^[129]

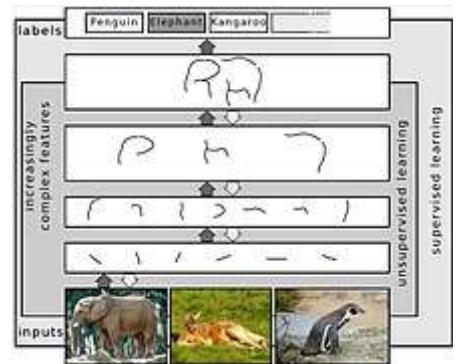
The main categories of networks are acyclic or feedforward neural networks (where the signal passes in only one direction) and recurrent neural networks (which allow feedback and short-term memories of previous input events). Among the most popular feedforward networks are perceptrons, multi-layer perceptrons and radial basis networks.^[130]

Deep learning

Deep learning^[132] uses several layers of neurons between the network's inputs and outputs. The multiple layers can progressively extract higher-level features from the raw input. For example, in image processing, lower layers may identify edges, while higher layers may identify the concepts relevant to a human such as digits or letters or faces.^[133] Deep learning has drastically improved the performance of programs in many important subfields of artificial intelligence, including computer vision, speech recognition, image classification^[134] and others.

Deep learning often uses convolutional neural networks for many or all of its layers. In a convolutional layer, each neuron receives input from only a restricted area of the previous layer called the neuron's receptive field. This can substantially reduce the number of weighted connections between neurons,^[135] and creates a hierarchy similar to the organization of the animal visual cortex.^[136]

In a recurrent neural network (RNN) the signal will propagate through a layer more than once;^[137] thus, an RNN is an example of deep learning.^[138] RNNs can be trained by gradient descent,^[139] however long-term gradients which are back-propagated can "vanish" (that is, they can tend to zero) or "explode" (that is, they can tend to infinity), known as the vanishing gradient problem.^[140] The long short term memory (LSTM) technique can prevent this in most cases.^[141]



Representing images on multiple layers of abstraction in deep learning^[131]

Specialized languages and hardware

Specialized languages for artificial intelligence have been developed, such as Lisp, Prolog, TensorFlow and many others. Hardware developed for AI includes AI accelerators and neuromorphic computing.

Applications

AI is relevant to any intellectual task.^[142] Modern artificial intelligence techniques are pervasive and are too numerous to list here.^[143] Frequently, when a technique reaches mainstream use, it is no longer considered artificial intelligence; this phenomenon is described as the AI effect.^[144]

In the 2010s, AI applications were at the heart of the most commercially successful areas of computing, and have become a ubiquitous feature of daily life. AI is used in search engines (such as Google Search), targeting online advertisements,^[145] recommendation systems (offered by Netflix, YouTube or Amazon), driving internet traffic,^{[146][147]} targeted advertising (AdSense, Facebook), virtual assistants (such as Siri or Alexa),^[148] autonomous vehicles (including drones

and self-driving cars), automatic translation (Microsoft Translator, Google Translate), face recognition (Apple's Face ID or Microsoft's DeepFace), image labeling (used by Facebook, Apple's iPhoto and TikTok) and spam filtering.

There are also thousands of successful AI applications used to solve problems for specific industries or institutions. A few examples are energy storage,^[149] deepfakes,^[150] medical diagnosis, military logistics, or supply chain management.

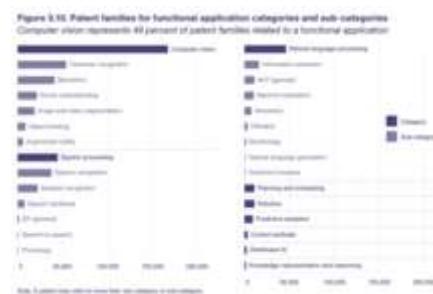
Game playing has been a test of AI's strength since the 1950s. Deep Blue became the first computer chess-playing system to beat a reigning world chess champion, Garry Kasparov, on 11 May 1997.^[151] In 2011, in a *Jeopardy!* quiz show exhibition match, IBM's question answering system, Watson, defeated the two greatest *Jeopardy!* champions, Brad Rutter and Ken Jennings, by a significant margin.^[152] In March 2016, AlphaGo won 4 out of 5 games of Go in a match with Go champion Lee Sedol, becoming the first computer Go-playing system to beat a professional Go player without handicaps.^[153] Other programs handle imperfect-information games; such as for poker at a superhuman level, Pluribus^[q] and Cepheus.^[155] DeepMind in the 2010s developed a "generalized artificial intelligence" that could learn many diverse Atari games on its own.^[156]

By 2020, Natural Language Processing systems such as the enormous GPT-3 (then by far the largest artificial neural network) were matching human performance on pre-existing benchmarks, albeit without the system attaining a commonsense understanding of the contents of the benchmarks.^[157] DeepMind's AlphaFold 2 (2020) demonstrated the ability to approximate, in hours rather than months, the 3D structure of a protein.^[158] Other applications predict the result of judicial decisions,^[159] create art (such as poetry or painting) and prove mathematical theorems.

In 2019, WIPO reported that AI was the most prolific emerging technology in terms of number of patent applications and granted patents, the Internet of things was estimated to be the largest in terms of market size. It was followed, again in market size, by big data technologies, robotics, AI, 3D printing and the fifth generation of mobile services (5G).^[160] Since AI emerged in the 1950s, 340,000 AI-related patent applications were filed by innovators and 1.6 million scientific papers have been published by researchers, with the majority of all AI-related patent filings published since 2013. Companies represent 26 out of the top 30 AI patent applicants, with universities or public research organizations accounting for the remaining four.^[161] The ratio of scientific papers to inventions has significantly decreased from 8:1 in 2010 to 3:1 in 2016, which is attributed to be indicative of a shift from theoretical research to the use of AI technologies in commercial products and services. Machine learning is the dominant AI technique disclosed in patents and is included in more than one-third of all identified inventions (134,777 machine learning patents filed for a total of 167,038 AI patents filed in 2016), with computer vision being the most popular functional application. AI-related patents not only disclose AI techniques and applications, they often also refer to an application field or industry. Twenty application fields were identified in 2016 and



For this project the AI had to learn the typical patterns in the colors and brushstrokes of Renaissance painter Raphael. The portrait shows the face of the actress Ornella Muti, "painted" by AI in the style of Raphael.



AI Patent families for functional application categories and sub-categories. Computer vision represents 49 percent of patent families related to a functional application in 2016.

included, in order of magnitude: tele and medical sciences (12 percent), interaction (11 percent). Other sectors included banking, computing and human-computer interaction (11 percent), transportation (15 percent), life science (10 percent), computing and human-computer interaction (11 percent). Other sectors included banking, entertainment, security, industry and manufacturing, agriculture, and networks (including social networks, smart cities and the Internet of things). IBM has the largest portfolio of AI patents with 8,290 patent applications, followed by Microsoft with 5,930 patent applications.^[161]

Philosophy

Defining artificial intelligence

Alan Turing wrote in 1950 "I propose to consider the question 'can machines think?'"^[162] He advised changing the question from whether a machine "thinks", to "whether or not it is possible for machinery to show intelligent behaviour".^[163] He devised the Turing test, which measures the ability of a machine to simulate human conversation.^[164] The only thing visible is the behavior of the machine, so it does not matter if the machine is conscious, or has a mind, or whether the intelligence is merely a "simulation" and not "the real thing".^[r]

Russell and Norvig are critical of the Turing test if it is used as the definition of artificial intelligence. "Aeronautical engineering texts," they wrote, "do not define the goal of their field as making 'machines that fly so exactly like pigeons that they can fool other pigeons.'"^[166] AI founder John McCarthy agreed, writing that "Artificial intelligence is not, by definition, simulation of human intelligence".^{[167][s]}

The intelligent agent paradigm^[169] defines intelligent behavior in general, without reference to human beings. An intelligent agent is a system that perceives its environment and takes actions that maximize its chances of success. Any system that has goal-directed behavior can be analyzed as an intelligent agent: something as simple as a thermostat, as complex as a human being, as well as large systems such as firms, biomes or nations. The intelligent agent paradigm became widely accepted during the 1990s, and currently serves as the definition of the field.^[a]

The intelligent agent paradigm has other advantages for AI. It provides a reliable and scientific way to test programs; researchers can directly compare or even combine different approaches to isolated problems, by asking which agent is best at maximizing a given "goal function". It also gives them a common language to communicate with other fields – such as mathematical optimization (which is defined in terms of "goals") or economics (which uses the same definition of a "rational agent").^[170]

Evaluating approaches to AI

No established unifying theory or paradigm has guided AI research for most of its history.^[t] The unprecedented success of statistical machine learning in the 2010s eclipsed all other approaches (so much so that some sources, especially in the business world, use the term "artificial intelligence" to mean "machine learning with neural networks"). This approach is mostly sub-symbolic, neat, soft and narrow (see below). Critics argue that these questions may have to be revisited by future generations of AI researchers.

Symbolic AI and its limits

Symbolic AI (or "GOFAI")^[172] simula' they solve puzzles, express legal rea.

vious reasoning that people use when they solve puzzles, express legal rea. They were highly successful at "intelligent" tasks such as algebra or IQ tests. In the 1960s, Newell and Simon proposed the physical symbol systems hypothesis: "A physical symbol system has the necessary and sufficient means of general intelligent action."^[173]

However, the symbolic approach failed dismally on many tasks that humans solve easily, such as learning, recognizing an object or commonsense reasoning. Moravec's paradox is the discovery that high-level "intelligent" tasks were easy for AI, but low level "instinctive" tasks were extremely difficult.^[174] Philosopher Hubert Dreyfus had argued since the 1960s that human expertise depends on unconscious instinct rather than conscious symbol manipulation, and on having a "feel" for the situation, rather than explicit symbolic knowledge.^[175] Although his arguments had been ridiculed and ignored when they were first presented, eventually, AI research came to agree.^{[u][48]}

The issue is not resolved: sub-symbolic reasoning can make many of the same inscrutable mistakes that human intuition does, such as algorithmic bias. Critics such as Noam Chomsky argue continuing research into symbolic AI will still be necessary to attain general intelligence,^{[177][178]} in part because sub-symbolic AI is a move away from explainable AI: it can be difficult or impossible to understand why a modern statistical AI program made a particular decision. The emerging field of neurosymbolic artificial intelligence attempts to bridge the two approaches.

Neat vs. scruffy

"Neats" hope that intelligent behavior is described using simple, elegant principles (such as logic, optimization, or neural networks). "Scruffies" expect that it necessarily requires solving a large number of unrelated problems (especially in areas like common sense reasoning). This issue was actively discussed in the 70s and 80s,^[179] but in the 1990s mathematical methods and solid scientific standards became the norm, a transition that Russell and Norvig termed "the victory of the neats".^[180]

Soft vs. hard computing

Finding a provably correct or optimal solution is intractable for many important problems.^[47] Soft computing is a set of techniques, including genetic algorithms, fuzzy logic and neural networks, that are tolerant of imprecision, uncertainty, partial truth and approximation. Soft computing was introduced in the late 80s and most successful AI programs in the 21st century are examples of soft computing with neural networks.

Narrow vs. general AI

AI researchers are divided as to whether to pursue the goals of artificial general intelligence and superintelligence (general AI) directly or to solve as many specific problems as possible (narrow AI) in hopes these solutions will lead indirectly to the field's long-term goals^{[181][182]} General intelligence is difficult to define and difficult to measure, and modern AI has had more verifiable successes by focussing on specific problems with specific solutions. The experimental sub-field of artificial general intelligence studies this area exclusively.

Machine consciousness, sentience and mind

The philosophy of mind does not know if machines can have a mind, consciousness and mental states, in the same sense. This issue considers the internal experiences of the machine, rather than its external behavior. Mainstream AI research considers this issue irrelevant because it does not affect the goals of the field. Stuart Russell and Peter Norvig observe that most AI researchers "don't care about the [philosophy of AI] – as long as the program works, they don't care whether you call it a simulation of intelligence or real intelligence."^[183] However, the question has become central to the philosophy of mind. It is also typically the central question at issue in artificial intelligence in fiction.

Consciousness

David Chalmers identified two problems in understanding the mind, which he named the "hard" and "easy" problems of consciousness.^[184] The easy problem is understanding how the brain processes signals, makes plans and controls behavior. The hard problem is explaining how this feels or why it should feel like anything at all. Human information processing is easy to explain, however, human subjective experience is difficult to explain. For example, it is easy to imagine a color-blind person who has learned to identify which objects in their field of view are red, but it is not clear what would be required for the person to *know what red looks like*.^[185]

Computationalism and functionalism

Computationalism is the position in the philosophy of mind that the human mind is an information processing system and that thinking is a form of computing. Computationalism argues that the relationship between mind and body is similar or identical to the relationship between software and hardware and thus may be a solution to the mind-body problem. This philosophical position was inspired by the work of AI researchers and cognitive scientists in the 1960s and was originally proposed by philosophers Jerry Fodor and Hilary Putnam.^[186]

Philosopher John Searle characterized this position as "strong AI": "The appropriately programmed computer with the right inputs and outputs would thereby have a mind in exactly the same sense human beings have minds."^[v] Searle counters this assertion with his Chinese room argument, which attempts to show that, even if a machine perfectly simulates human behavior, there is still no reason to suppose it also has a mind.^[189]

Robot rights

If a machine has a mind and subjective experience, then it may also have sentience (the ability to feel), and if so, then it could also suffer, and thus it would be entitled to certain rights.^[190] Any hypothetical robot rights would lie on a spectrum with animal rights and human rights.^[191] This issue has been considered in fiction for centuries,^[192] and is now being considered by, for example, California's Institute for the Future; however, critics argue that the discussion is premature.^[193]

Future

Superintelligence

A superintelligence, hyperintelligence, or superhuman intelligence, is a hypothetical agent that would possess intelligence far surpassing that of the brightest and most gifted human mind. Superintelligence may also refer to the form or degree of intelligence possessed by such an

agent.^[182]



If research into artificial general intelligence produced sufficiently intelligent software, it might be able to reprogram and improve itself. The improved software would be even better at improving itself, leading to recursive self-improvement.^[194] Its intelligence would increase exponentially in an intelligence explosion and could dramatically surpass humans. Science fiction writer Vernor Vinge named this scenario the "singularity".^[195] Because it is difficult or impossible to know the limits of intelligence or the capabilities of superintelligent machines, the technological singularity is an occurrence beyond which events are unpredictable or even unfathomable.^[196]

Robot designer Hans Moravec, cyberneticist Kevin Warwick, and inventor Ray Kurzweil have predicted that humans and machines will merge in the future into cyborgs that are more capable and powerful than either. This idea, called transhumanism, has roots in Aldous Huxley and Robert Ettinger.^[197]

Edward Fredkin argues that "artificial intelligence is the next stage in evolution", an idea first proposed by Samuel Butler's "Darwin among the Machines" as far back as 1863, and expanded upon by George Dyson in his book of the same name in 1998.^[198]

Risks

Technological unemployment

In the past technology has tended to increase rather than reduce total employment, but economists acknowledge that "we're in uncharted territory" with AI.^[199] A survey of economists showed disagreement about whether the increasing use of robots and AI will cause a substantial increase in long-term unemployment, but they generally agree that it could be a net benefit if productivity gains are redistributed.^[200] Subjective estimates of the risk vary widely; for example, Michael Osborne and Carl Benedikt Frey estimate 47% of U.S. jobs are at "high risk" of potential automation, while an OECD report classifies only 9% of U.S. jobs as "high risk".^{[w][202]}

Unlike previous waves of automation, many middle-class jobs may be eliminated by artificial intelligence; The Economist states that "the worry that AI could do to white-collar jobs what steam power did to blue-collar ones during the Industrial Revolution" is "worth taking seriously".^[203] Jobs at extreme risk range from paralegals to fast food cooks, while job demand is likely to increase for care-related professions ranging from personal healthcare to the clergy.^[204]

Bad actors and weaponized AI

AI provides a number of tools that are particularly useful for authoritarian governments: smart spyware, face recognition and voice recognition allow widespread surveillance; such surveillance allows machine learning to classify potential enemies of the state and can prevent them from hiding; recommendation systems can precisely target propaganda and misinformation for maximum effect; deepfakes aid in producing misinformation; advanced AI can make centralized decision making more competitive with liberal and decentralized systems such as markets.^[205]

Terrorists, criminals and rogue states may use other forms of weaponized AI such as advanced digital warfare and lethal autonomous weapons. By 2015, over fifty countries were reported to be researching battlefield robots.^[206]

Machine-learning AI is also able to design tens of thousands of toxic molecules in a matter of hours.^[207]

Algorithmic bias



AI programs can become biased after learning from real-world data. It is not typically introduced by the system designers but is learned by the program, and thus the programmers are often unaware that the bias exists.^[208] Bias can be inadvertently introduced by the way training data is selected.^[209] It can also emerge from correlations: AI is used to classify individuals into groups and then make predictions assuming that the individual will resemble other members of the group. In some cases, this assumption may be unfair.^[210] An example of this is COMPAS, a commercial program widely used by U.S. courts to assess the likelihood of a defendant becoming a recidivist. ProPublica claims that the COMPAS-assigned recidivism risk level of black defendants is far more likely to be overestimated than that of white defendants, despite the fact that the program was not told the races of the defendants.^[211] Other examples where algorithmic bias can lead to unfair outcomes are when AI is used for credit rating or hiring.

At its 2022 Conference on Fairness, Accountability, and Transparency (ACM FAccT 2022) the Association for Computing Machinery, in Seoul, South Korea, presented and published findings recommending that until AI and robotics systems are demonstrated to be free of bias mistakes, they are unsafe and the use of self-learning neural networks trained on vast, unregulated sources of flawed internet data should be curtailed.^[212]

Existential risk

Superintelligent AI may be able to improve itself to the point that humans could not control it. This could, as physicist Stephen Hawking puts it, "spell the end of the human race".^[213] Philosopher Nick Bostrom argues that sufficiently intelligent AI, if it chooses actions based on achieving some goal, will exhibit convergent behavior such as acquiring resources or protecting itself from being shut down. If this AI's goals do not fully reflect humanity's, it might need to harm humanity to acquire more resources or prevent itself from being shut down, ultimately to better achieve its goal. He concludes that AI poses a risk to mankind, however humble or "friendly" its stated goals might be.^[214] Political scientist Charles T. Rubin argues that "any sufficiently advanced benevolence may be indistinguishable from malevolence." Humans should not assume machines or robots would treat us favorably because there is no *a priori* reason to believe that they would share our system of morality.^[215]

The opinion of experts and industry insiders is mixed, with sizable fractions both concerned and unconcerned by risk from eventual superhumanly-capable AI.^[216] Stephen Hawking, Microsoft founder Bill Gates, history professor Yuval Noah Harari, and SpaceX founder Elon Musk have all expressed serious misgivings about the future of AI.^[217] Prominent tech titans including Peter Thiel (Amazon Web Services) and Musk have committed more than \$1 billion to nonprofit companies that champion responsible AI development, such as OpenAI and the Future of Life Institute.^[218] Mark Zuckerberg (CEO, Facebook) has said that artificial intelligence is helpful in its current form and will continue to assist humans.^[219] Other experts argue is that the risks are far enough in the future to not be worth researching, or that humans will be valuable from the perspective of a superintelligent machine.^[220] Rodney Brooks, in particular, has said that "malevolent" AI is still centuries away.^[x]

Other issues: Copyright

AI's decisions making abilities raises the questions of legal responsibility and copyright status of created works. These issues are being refined in various jurisdictions.^[222]

Ethical machines

Friendly AI are machines that have  beginning to minimize risks and to make choices that benefit humans  who coined the term, argues that developing friendly AI should be a higher research priority: it may require a large investment and it must be completed before AI becomes an existential risk.^[223]

Machines with intelligence have the potential to use their intelligence to make ethical decisions. The field of machine ethics provides machines with ethical principles and procedures for resolving ethical dilemmas.^[224] Machine ethics is also called machine morality, computational ethics or computational morality,^[224] and was founded at an AAAI symposium in 2005.^[225]

Other approaches include Wendell Wallach's "artificial moral agents"^[226] and Stuart J. Russell's three principles for developing provably beneficial machines.^[227]

Regulation

The regulation of artificial intelligence is the development of public sector policies and laws for promoting and regulating artificial intelligence (AI); it is therefore related to the broader regulation of algorithms.^[228] The regulatory and policy landscape for AI is an emerging issue in jurisdictions globally.^[229] Between 2016 and 2020, more than 30 countries adopted dedicated strategies for AI.^[44] Most EU member states had released national AI strategies, as had Canada, China, India, Japan, Mauritius, the Russian Federation, Saudi Arabia, United Arab Emirates, US and Vietnam. Others were in the process of elaborating their own AI strategy, including Bangladesh, Malaysia and Tunisia.^[44] The Global Partnership on Artificial Intelligence was launched in June 2020, stating a need for AI to be developed in accordance with human rights and democratic values, to ensure public confidence and trust in the technology.^[44] Henry Kissinger, Eric Schmidt, and Daniel Huttenlocher published a joint statement in November 2021 calling for a government commission to regulate AI.^[230]

In fiction

Thought-capable artificial beings have appeared as storytelling devices since antiquity,^[15] and have been a persistent theme in science fiction.^[17]

A common trope in these works began with Mary Shelley's Frankenstein, where a human creation becomes a threat to its masters. This includes such works as Arthur C. Clarke's and Stanley Kubrick's 2001: A Space Odyssey (both 1968), with HAL 9000, the murderous computer in charge of the Discovery One spaceship, as well as The Terminator (1984) and The Matrix (1999). In contrast, the rare loyal robots such as Gort from The Day the Earth Stood Still (1951) and Bishop from Aliens (1986) are less prominent in popular culture.^[231]



The word "robot" itself was coined by Karel Čapek in his 1921 play R.U.R., the title standing for "Rossum's Universal Robots".

Isaac Asimov introduced the Three Laws of Robotics in many books and stories, most notably the "Multivac" series about a super-intelligent computer of the same name. Asimov's laws are often brought up during lay discussions of machine ethics;^[232] while almost all artificial intelligence researchers are familiar with Asimov's laws through popular culture, they generally consider the laws useless for many reasons, one of which is their ambiguity.^[233]

Transhumanism (the merging of humans and machines) is explored in the manga *Ghost in the Shell* and the science-fiction series *Dune*.

Several works use AI to force us to consider the question of what makes us human, showing us artificial beings that have thus to suffer. This appears in Karel Čapek's *R.U.R.*, the films *A.I. Artificial Intelligence* and *Ex Machina*, as well as the novel *Do Androids Dream of Electric Sheep?*, by Philip K. Dick. Dick considers the idea that our understanding of human subjectivity is altered by technology created with artificial intelligence.^[234]

Scientific diplomacy

Warfare

As technology and research evolve and the world enters the third revolution of warfare following gunpowder and nuclear weapons, the artificial intelligence arms race ensues between the United States, China, and Russia, three countries with the world's top five highest military budgets.^[235] Intentions of being a world leader in AI research by 2030^[236] have been declared by China's leader Xi Jinping, and President Putin of Russia has stated that "Whoever becomes the leader in this sphere will become the ruler of the world".^[237] If Russia were to become the leader in AI research, President Putin has stated Russia's intent to share some of their research with the world so as to not monopolize the field,^[237] similar to their current sharing of nuclear technologies, maintaining science diplomacy relations. The United States, China, and Russia, are some examples of countries that have taken their stances toward military artificial intelligence since as early as 2014, having established military programs to develop cyber weapons, control lethal autonomous weapons, and drones that can be used for surveillance.

Russo-Ukrainian War

President Putin announced that artificial intelligence is the future for all mankind^[237] and recognizes the power and opportunities that the development and deployment of lethal autonomous weapons AI technology can hold in warfare and homeland security, as well as its threats. President Putin's prediction that future wars will be fought using AI has started to come to fruition to an extent after Russia invaded Ukraine on 24 February 2022. The Ukrainian military is making use of the Turkish Bayraktar TB2-drones^[238] that still require human operation to deploy laser-guided bombs but can take off, land, and cruise autonomously. Ukraine has also been using Switchblade drones supplied by the US and receiving information gathering by the United States's own surveillance operations regarding battlefield intelligence and national security about Russia.^[239] Similarly, Russia can use AI to help analyze battlefield data from surveillance footage taken by drones. Reports and images show that Russia's military has deployed KUB- BLA suicide drones^[240] into Ukraine, with speculations of intentions to assassinate Ukrainian President Volodymyr Zelenskyy.

Warfare regulations

As research in the AI realm progresses, there is pushback about the use of AI from the Campaign to Stop Killer Robots and world technology leaders have sent a petition^[241] to the United Nations calling for new regulations on the development and use of AI technologies in 2017, including a ban on the use of lethal autonomous weapons due to ethical concerns for innocent civilian populations.

Cybersecurity

With the ever evolving cyber-attacks and more effective response by rising actors, AI can be used for threat detection and this tool, some challenges are also presented such as privacy, informed consent, and responsible use.^[242] According to CISA, the cyberspace is difficult to secure for the following factors: the ability of malicious actors to operate from anywhere in the world, the linkages between cyberspace and physical systems, and the difficulty of reducing vulnerabilities and consequences in complex cyber networks.^[243] With the increased technological advances of the world, the risk for wide scale consequential events rises. Paradoxically, the ability to protect information and create a line of communication between the scientific and diplomatic community thrives. The role of cybersecurity in diplomacy has become increasingly relevant, creating the term of cyber diplomacy – which is not uniformly defined and not synonymous with cyber defence.^[244] Many nations have developed unique approaches to scientific diplomacy in cyberspace.

Czech Republic's approach

Dating back to 2011, when the Czech National Security Authority (NSA) was appointed as the national authority for the cyber agenda. The role of cyber diplomacy strengthened in 2017 when the Czech Ministry of Foreign Affairs (MFA) detected a serious cyber campaign directed against its own computer networks.^[245] In 2016, three cyber diplomats were deployed to Washington, D.C., Brussels and Tel Aviv, with the goal of establishing active international cooperation focused on engagement with the EU and NATO. The main agenda for these scientific diplomacy efforts is to bolster research on artificial intelligence and how it can be utilized in cybersecurity research, development, and overall consumer trust.^[246] CzechInvest is a key stakeholder in scientific diplomacy and cybersecurity. For example, in September 2018, they organized a mission to Canada in September 2018 with a special focus on artificial intelligence. The main goal of this particular mission was a promotional effort on behalf of Prague, attempting to establish it as a future knowledge hub for the industry for interested Canadian firms.^[247]

Germany's approach

Cybersecurity is recognized as a governmental task, dividing into three ministries of responsibility: the Federal Ministry of the Interior, the Federal Ministry of Defence, and the Federal Foreign Office.^[248] These distinctions promoted the creation of various institutions, such as The German National Office for Information Security, The National Cyberdefence Centre, The German National Cyber Security Council, and The Cyber and Information Domain Service.^[246] In 2018, a new strategy for artificial intelligence was established by the German government, with the creation of a German-French virtual research and innovation network,^[249] holding opportunity for research expansion into cybersecurity.

European Union's approach

The adoption of *The Cybersecurity Strategy of the European Union – An Open, Safe and Secure Cyberspace* document in 2013 by the European commission^[246] pushed forth cybersecurity efforts integrated with scientific diplomacy and artificial intelligence. Efforts are strong, as the EU funds various programs and institutions in the effort to bring science to diplomacy and bring diplomacy to science. Some examples are the cyber security programme Competence Research Innovation (CONCORDIA), which brings together 14 member states,^[250] and Cybersecurity for Europe (CSE), which brings together 43 partners involving 20 member states.^[251] In addition, The European Network of Cybersecurity Centres and Competence Hub for Innovation and Operations (ECHO) gathers 30 partners with 15 member states^[252] and SPARTA gathers 44 partners

involving 14 member states.^[253] The overall goals of the EU, to innovate cybersecurity for defense and protect integrated cyberspace among many nations, and further contribute to the security of artificial intelligence.^[246]

Russo-Ukrainian War

With the 2022 invasion of Ukraine, there has been a rise in malicious cyber activity against the United States,^[254] Ukraine, and Russia. A prominent and rare documented use of artificial intelligence in conflict is on behalf of Ukraine, using facial recognition software to uncover Russian assailants and identify Ukrainians killed in the ongoing war.^[255] Though these governmental figures are not primarily focused on scientific and cyber diplomacy, other institutions are commenting on the use of artificial intelligence in cybersecurity with that focus. For example, Georgetown University's Center for Security and Emerging Technology (CSET) has the Cyber-AI Project, with one goal being to attract policymakers' attention to the growing body of academic research, which exposes the exploitative consequences of AI and machine-learning (ML) algorithms.^[256] This vulnerability can be a plausible explanation as to why Russia is not engaging in the use of AI in conflict per, Andrew Lohn, a senior fellow at CSET. In addition to use on the battlefield, AI is being used by the Pentagon to analyze data from the war, analyzing to strengthen cybersecurity and warfare intelligence for the United States.^{[239][257]}

Election security

As artificial intelligence grows and the overwhelming amount of news portrayed through cyberspace expands, it is becoming extremely overwhelming for a voter to know what to believe. There are many intelligent codes, referred to as bots, written to portray people on social media with the goal of spreading miss information.^[258] The 2016 US election is a victim of such actions. During the Hillary Clinton and Donald Trump campaign, artificial intelligent bots from Russia were spreading misinformation about the candidates in order to help the Trump campaign.^[259] Analysts concluded that approximately 19% of Twitter tweets centered around the 2016 election were detected to come from bots.^[259] YouTube in recent years has been used to spread political information as well. Although there is no proof that the platform attempts to manipulate its viewers opinions, Youtubes AI algorithm recommends videos of similar variety.^[260] If a person begins to research right wing political podcasts, then YouTube's algorithm will recommend more right wing videos.^[261] The uprising in a program called Deepfake, a software used to replicate someone's face and words, has also shown its potential threat. In 2018 a Deepfake video of Barack Obama was released saying words he claims to have never said.^[262] While in a national election a Deepfake will quickly be debunked, the software has the capability to heavily sway a smaller local election. This tool holds a lot of potential for spreading misinformation and is monitored with great attention.^[263] Although it may be seen as a tool used for harm, AI can help enhance election campaigns as well. AI bots can be programed to target articles with known misinformation. The bots can then indicate what is being misinformed to help shine light on the truth. AI can also be used to inform a person where each parts stands on a certain topic such as healthcare or climate change.^[264] The political leaders of a nation have heavy sway on international affairs. Thus, a political leader with a lack of interest for international collaborative scientific advancement can have a negative impact in the scientific diplomacy of that nation^[265]

Future of work

Facial recognition

The use of artificial intelligence (AI) has come part of everyday life. It is used every day in facial recognition software. This means of security for many companies in the form of a biometric authentication. This means of authentication allows even the most official organizations such as the United States Internal Revenue Service to verify a person's identity [266] via a database generated from machine learning. As of the year 2022, the United States IRS requires those who do not undergo a live interview with an agent to complete a biometric verification of their identity via ID.me's facial recognition tool. [266]

AI and school

In Japan and South Korea, artificial intelligence software is used in the instruction of English language via the company Riid. [267] Riid is a Korean education company working alongside Japan to give students the means to learn and use their English communication skills via engaging with artificial intelligence in a live chat. [267] Riid is not the only company to do this. American company Duolingo is well known for their automated teaching of 41 languages. Babbel, a German language learning program, also uses artificial intelligence in its teaching automation, allowing for European students to learn vital communication skills needed in social, economic, and diplomatic settings. Artificial intelligence will also automate the routine tasks that teachers need to do such as grading, taking attendance, and handling routine student inquiries. [268] This enables the teacher to carry on with the complexities of teaching that an automated machine cannot handle. These include creating exams, explaining complex material in a way that will benefit students individually and handling unique questions from students.

AI and medicine

Unlike the human brain, which possess generalized intelligence, the specialized intelligence of AI can serve as a means of support to physicians internationally. The medical field has a diverse and profound amount of data in which AI can employ to generate a predictive diagnosis. Researchers at an Oxford hospital have developed artificial intelligence that can diagnose heart scans for heart disease and cancer. [269] This artificial intelligence can pick up diminutive details in the scans that doctors may miss. As such, artificial intelligence in medicine will better the industry, giving doctors the means to precisely diagnose their patients using the tools available. The artificial intelligence algorithms will also be used to further improve diagnosis over time, via an application of machine learning called precision medicine. [270] Furthermore, the narrow application of artificial intelligence can use "deep learning" in order to improve medical image analysis. In radiology imaging, AI uses deep learning algorithms to identify potentially cancerous lesions which is an important process assisting in early diagnosis. [271]

AI in business

Data analysis is a fundamental property of artificial intelligence that enables it to be used in every facet of life from search results to the way people buy product. According to NewVantage Partners, [272] over 90% of top businesses have ongoing investments in artificial intelligence. According to IBM, one of the world's leaders in technology, 45% of respondents from companies with over 1,000 employees have adopted AI. [273] Recent data shows that the business market [274] for artificial intelligence during the year 2020 was valued at \$51.08 billion. The business market for artificial intelligence is projected to be over \$640.3 billion by the year 2028. [274] To prevent harm, AI-deploying organizations need to play a central role in creating and deploying trustworthy AI in line with the principles of trustworthy AI, [275] and take accountability to mitigate the risks. [276]

Business and diplomacy



With the exponential surge of artificial technology and communication, the distribution of one's ideals and values has been evident in daily life. Digital information is spread via communication apps such as Whatsapp, Facebook/Meta, Snapchat, Instagram and Twitter. However, it is known that these sites relay specific information corresponding to data analysis. If a right-winged individual were to do a google search, Google's algorithms would target that individual and relay data pertinent to that target audience. US President Bill Clinton noted in 2000: "In the new century, liberty will spread by cell phone and cable modem. [...] We know how much the Internet has changed America, and we are already an open society."^[277] However, when the private sector uses artificial intelligence to gather data, a shift in power from the state to the private sector may be seen. This shift in power, specifically in large technological corporations, could profoundly change how diplomacy functions in society. The rise in digital technology and usage of artificial technology enabled the private sector to gather immense data on the public, which is then further categorized by race, location, age, gender, etc.^[278] *The New York Times* calculates that "the ten largest tech firms, which have become gatekeepers in commerce, finance, entertainment and communications, now have a combined market capitalization of more than \$10 trillion. In gross domestic product terms, that would rank them as the world's third-largest economy."^[279] Beyond the general lobbying of congressmen/congresswomen, companies such as Facebook/Meta or Google use collected data in order to reach their intended audiences with targeted information.^[279]

AI and foreign policy

Multiple nations around the globe employ artificial intelligence to assist with their foreign policy decisions. The Chinese Department of External Security Affairs – under the Ministry of Foreign Affairs – uses AI to review almost all its foreign investment projects for risk mitigation.^[280] The government of China plans to utilize artificial intelligence in its \$900 billion global infrastructure development plan, called the "Belt and Road Initiative" for political, economic, and environmental risk alleviation.^[281]

Over 200 applications of artificial intelligence are being used by over 46 United Nations agencies, in sectors ranging from health care dealing with issues such as combating COVID-19 to smart agriculture, to assist the UN in political and diplomatic relations.^[282] One example is the use of AI by the UN Global Pulse program to model the effect of the spread of COVID-19 on internally displaced people (IDP) and refugee settlements to assist them in creating an appropriate global health policy.^{[283][284]}

Novel AI tools such as remote sensing can also be employed by diplomats for collecting and analyzing data and near-real-time tracking of objects such as troop or refugee movements along borders in violent conflict zones.^{[283][285]}

Artificial intelligence can be used to mitigate vital cross-national diplomatic talks to prevent translation errors caused by human translators.^[286] A major example is the 2021 Anchorage meetings held between US and China aimed at stabilizing foreign relations, only for it to have the opposite effect, increasing tension and aggressiveness between the two nations, due to translation errors caused by human translators.^[287] In the meeting, when United States National Security Advisor to President Joe Biden, Jacob Jeremiah Sullivan stated, "We do not seek conflict, but we welcome stiff competition and we will always stand up for our principles, for our people, and for our friends", it was mistranslated into Chinese as "we will face competition between us, and will present our stance in a very clear manner", adding an aggressive tone to the speech.^[287] AI's

ability for fast and efficient natural language processing makes it an important tool for communication and real-time translation and communication between nations and prevents unintended mistranslation.^[280]

ing and real-time translation and communication between nations and

See also

- [A.I. Rising](#)
- [AI control problem](#)
- [Artificial intelligence arms race](#)
- [Behavior selection algorithm](#)
- [Business process automation](#)
- [Case-based reasoning](#)
- [Emergent algorithm](#)
- [Female gendering of AI technologies](#)
- [Glossary of artificial intelligence](#)
- [Robotic process automation](#)
- [Synthetic intelligence](#)
- [Universal basic income](#)
- [Weak AI](#)

Explanatory notes

- a. Definition of AI as the study of [intelligent agents](#), drawn from leading AI textbooks.
 - [Poole, Mackworth & Goebel \(1998, p. 1 \(<http://people.cs.ubc.ca/~poole/ci/ch1.pdf>\)\)](#), which provides the version that is used in this article. These authors use the term "computational intelligence" as a synonym for artificial intelligence.
 - [Russell & Norvig \(2003, p. 55\)](#) (who prefer the term "rational agent") and write "The whole-agent view is now widely accepted in the field".
 - [Nilsson \(1998\)](#)
 - [Legg & Hutter \(2007\)](#)
- b. [Stuart Russell](#) and [Peter Norvig](#) characterize this definition as "thinking humanly" and reject it in favor of "acting rationally".^[1]
- c. This list of intelligent traits is based on the topics covered by the major AI textbooks, including: [Russell & Norvig \(2003\)](#), [Luger & Stubblefield \(2004\)](#), [Poole, Mackworth & Goebel \(1998\)](#) and [Nilsson \(1998\)](#)
- d. This statement comes from the proposal for the [Dartmouth workshop](#) of 1956, which reads: "Every aspect of learning or any other feature of intelligence can be so precisely described that a machine can be made to simulate it."^[13]
- e. Russel and Norvig note in the textbook Artificial Intelligence: A Modern Approach (4th ed.), section 1.5: "In the longer term, we face the difficult problem of controlling superintelligent AI systems that may evolve in unpredictable ways." while referring to computer scientists, philosophers, and technologists.
- f. Daniel Crevier wrote "the conference is generally recognized as the official birthdate of the new science."^[23] [Russell](#) and [Norvig](#) call the conference "the birth of artificial intelligence."^[24]
- g. [Russell](#) and [Norvig](#) wrote "for the next 20 years the field would be dominated by these people and their students."^[24]
- h. [Russell](#) and [Norvig](#) wrote "it was astonishing whenever a computer did anything kind of smartish".^[26]

- i. The programs described are Arthur Bobrow's STUDENT, Newell and Simon's General Problem Solver program for the IBM 701, Daniel Bobrow's STUDENT, Newell and Hartley Winograd's SHRDLU.
- j. Embodied approaches to AI^[36] were championed by Hans Moravec^[37] and Rodney Brooks^[38] and went by many names: Nouvelle AI,^[38] Developmental robotics,^[39] situated AI, behavior-based AI as well as others. A similar movement in cognitive science was the embodied mind thesis.
- k. Clark wrote: "After a half-decade of quiet breakthroughs in artificial intelligence, 2015 has been a landmark year. Computers are smarter and learning faster than ever."^[10]
- l. Alan Turing discussed the centrality of learning as early as 1950, in his classic paper "Computing Machinery and Intelligence".^[66] In 1956, at the original Dartmouth AI summer conference, Ray Solomonoff wrote a report on unsupervised probabilistic machine learning: "An Inductive Inference Machine".^[67]
- m. This is a form of Tom Mitchell's widely quoted definition of machine learning: "A computer program is set to learn from an experience E with respect to some task T and some performance measure P if its performance on T as measured by P improves with experience E ".^[68]
- n. Alan Turing suggested in "Computing Machinery and Intelligence" that a "thinking machine" would need to be educated like a child.^[66] Developmental robotics is a modern version of the idea.^[39]
- o. Compared with symbolic logic, formal Bayesian inference is computationally expensive. For inference to be tractable, most observations must be conditionally independent of one another. AdSense uses a Bayesian network with over 300 million edges to learn which ads to serve.^[110]
- p. Expectation-maximization, one of the most popular algorithms in machine learning, allows clustering in the presence of unknown latent variables.^[112]
- q. The Smithsonian reports: "Pluribus has bested poker pros in a series of six-player no-limit Texas Hold'em games, reaching a milestone in artificial intelligence research. It is the first bot to beat humans in a complex multiplayer competition."^[154]
- r. Russell and Norvig characterize this as the difference between "acting" and "thinking".^[165]
- s. The distinction between "acting humanly" and "acting rationally" is due to Russell and Norvig.^[165] Pamela McCorduck wrote in 2004 that there are "two major branches of artificial intelligence: one aimed at producing intelligent behavior regardless of how it was accomplished, and the other aimed at modeling intelligent processes found in nature, particularly human ones."^[168]
- t. Nils Nilsson wrote in 1983: "Simply put, there is wide disagreement in the field about what AI is all about."^[171]
- u. Daniel Crevier wrote that "time has proven the accuracy and perceptiveness of some of Dreyfus's comments. Had he formulated them less aggressively, constructive actions they suggested might have been taken much earlier."^[176]
- v. Searle presented this definition of "Strong AI" in 1999.^[187] Searle's original formulation was "The appropriately programmed computer really is a mind, in the sense that computers given the right programs can be literally said to understand and have other cognitive states."^[188] Strong AI is defined similarly by Russell and Norvig: "The assertion that machines could possibly act intelligently (or, perhaps better, act as if they were intelligent) is called the 'weak AI' hypothesis by philosophers, and the assertion that machines that do so are actually thinking (as opposed to simulating thinking) is called the 'strong AI' hypothesis."^[183]
- w. See table 4; 9% is both the OECD average and the US average.^[201]
- x. Rodney Brooks writes, "I think it is a mistake to be worrying about us developing malevolent AI anytime in the next few hundred years. I think the worry stems from a fundamental error in not distinguishing the difference between the very real recent advances in a particular aspect of AI and the enormity and complexity of building sentient volitional intelligence."^[221]

Citations



1. [Russell & Norvig \(2021\)](#), p. 2.
2. [Google \(2016\)](#).
3. [McCorduck \(2004\)](#), p. 204.
4. [Ashok83 \(2019\)](#).
5. [Schank \(1991\)](#), p. 38.
6. [Crevier \(1993\)](#), p. 109.
7. Funding initiatives in the early 80s: [Fifth Generation Project \(Japan\)](#), [Alvey \(UK\)](#), [Microelectronics and Computer Technology Corporation \(US\)](#), [Strategic Computing Initiative \(US\)](#):
 - [McCorduck \(2004\)](#), pp. 426–441
 - [Crevier \(1993\)](#), pp. 161–162, 197–203, 211, 240
 - [Russell & Norvig \(2003\)](#), p. 24
 - [NRC \(1999\)](#), pp. 210–211
 - [Newquist \(1994\)](#), pp. 235–248
8. First AI Winter, [Lighthill report](#), [Mansfield Amendment](#)
 - [Crevier \(1993\)](#), pp. 115–117
 - [Russell & Norvig \(2003\)](#), p. 22
 - [NRC \(1999\)](#), pp. 212–213
 - [Howe \(1994\)](#)
 - [Newquist \(1994\)](#), pp. 189–201
9. Second AI Winter:
 - [McCorduck \(2004\)](#), pp. 430–435
 - [Crevier \(1993\)](#), pp. 209–210
 - [NRC \(1999\)](#), pp. 214–216
 - [Newquist \(1994\)](#), pp. 301–318
10. [Clark \(2015b\)](#).
11. AI widely used in late 1990s:
 - [Russell & Norvig \(2003\)](#), p. 28
 - [Kurzweil \(2005\)](#), p. 265
 - [NRC \(1999\)](#), pp. 216–222
 - [Newquist \(1994\)](#), pp. 189–201
12. [Pennachin & Goertzel \(2007\)](#); [Roberts \(2016\)](#)
13. [McCarthy et al. \(1955\)](#).
14. [Newquist \(1994\)](#), pp. 45–53.
15. AI in myth:
 - [McCorduck \(2004\)](#), pp. 4–5
 - [Russell & Norvig \(2003\)](#), p. 939
16. [McCorduck \(2004\)](#), pp. 17–25.
17. [McCorduck \(2004\)](#), pp. 340–400.
18. [Berlinski \(2000\)](#).
19. AI's immediate precursors:
 - [McCorduck \(2004\)](#), pp. 51–107
 - [Crevier \(1993\)](#), pp. 27–32
 - [Russell & Norvig \(2003\)](#), pp. 15, 940
 - [Moravec \(1988\)](#), p. 3
20. [Russell & Norvig \(2009\)](#), p. 16.
21. [Manyika 2022](#), p. 9.
22. [Manyika 2022](#), p. 10.
23. [Crevier \(1993\)](#), pp. 47–49.
24. [Russell & Norvig \(2003\)](#), p. 17.
25. Dartmouth workshop:
 - [Russell & Norvig \(2003\)](#), p. 17
 - [McCorduck \(2004\)](#), pp. 111–136
 - [NRC \(1999\)](#), pp. 200–201

The proposal:

 - [McCarthy et al. \(1955\)](#)
26. [Russell & Norvig \(2003\)](#), p. 18.
27. Successful [Symbolic AI](#) programs:
 - [McCorduck \(2004\)](#), pp. 243–252
 - [Crevier \(1993\)](#), pp. 52–107
 - [Moravec \(1988\)](#), p. 9
 - [Russell & Norvig \(2003\)](#), pp. 18–21
28. AI heavily funded in 1960s:
 - [McCorduck \(2004\)](#), p. 131
 - [Crevier \(1993\)](#), pp. 51, 64–65
 - [NRC \(1999\)](#), pp. 204–205
29. [Howe \(1994\)](#).
30. [Newquist \(1994\)](#), pp. 86–86.
31. [Simon \(1965\)](#), p. 96) quoted in [Crevier \(1993\)](#), p. 109)
32. [Minsky \(1967\)](#), p. 2) quoted in [Crevier \(1993\)](#), p. 109)
33. [Lighthill \(1973\)](#).

34. Expert systems:

- [Russell & Norvig \(2003, pp. 22–23\)](#)
- [Luger & Stubblefield \(2004, pp. 227–331\)](#)
- [Nilsson \(1998, chpt. 17.4\)](#)
- [McCorduck \(2004, pp. 327–335, 434–435\)](#)
- [Crevier \(1993, pp. 145–62, 197–203\)](#)
- [Newquist \(1994, pp. 155–183\)](#)

35. [Nilsson \(1998\)](#), p. 7.

36. [McCorduck \(2004\)](#), pp. 454–462.

37. [Moravec \(1988\)](#).

38. [Brooks \(1990\)](#).

39. Developmental robotics:

- [Weng et al. \(2001\)](#)
- [Lungarella et al. \(2003\)](#)
- [Asada et al. \(2009\)](#)
- [Oudeyer \(2010\)](#)

40. Revival of connectionism:

- [Crevier \(1993, pp. 214–215\)](#)
- [Russell & Norvig \(2003, p. 25\)](#)

41. Formal and narrow methods adopted in the 1990s:

- [Russell & Norvig \(2003, pp. 25–26\)](#)
- [McCorduck \(2004, pp. 486–487\)](#)

42. [McKinsey \(2018\)](#).

43. [MIT Sloan Management Review \(2018\)](#); [Lorica \(2017\)](#)

44. [UNESCO \(2021\)](#).

45. Problem solving, puzzle solving, game playing and deduction:

- [Russell & Norvig \(2003, chpt. 3–9\)](#)
- [Poole, Mackworth & Goebel \(1998, chpts. 2,3,7,9\)](#)
- [Luger & Stubblefield \(2004, chpt. 3,4,6,8\)](#)
- [Nilsson \(1998, chpt. 7–12\)](#)

46. Uncertain reasoning:

- [Russell & Norvig \(2003, pp. 452–644\)](#)
- [Poole, Mackworth & Goebel \(1998, pp. 345–395\)](#)
- [Luger & Stubblefield \(2004, pp. 333–381\)](#)
- [Nilsson \(1998, chpt. 19\)](#)

ability and efficiency and the
editorial explosion:

- [Russell & Norvig \(2003, pp. 9, 21–22\)](#)

48. Psychological evidence of the prevalence of sub-symbolic reasoning and knowledge:

- [Kahneman \(2011\)](#)
- [Wason & Shapiro \(1966\)](#)
- [Kahneman, Slovic & Tversky \(1982\)](#)
- [Dreyfus & Dreyfus \(1986\)](#)

49. Knowledge representation and knowledge engineering:

- [Russell & Norvig \(2003, pp. 260–266, 320–363\)](#)
- [Poole, Mackworth & Goebel \(1998, pp. 23–46, 69–81, 169–233, 235–277, 281–298, 319–345\)](#)
- [Luger & Stubblefield \(2004, pp. 227–243\)](#),
- [Nilsson \(1998, chpt. 17.1–17.4, 18\)](#)

50. [Russell & Norvig \(2003\)](#), pp. 320–328.

51. Representing categories and relations:
Semantic networks, description logics,
inheritance (including frames and scripts):

- [Russell & Norvig \(2003, pp. 349–354\)](#),
- [Poole, Mackworth & Goebel \(1998, pp. 174–177\)](#),
- [Luger & Stubblefield \(2004, pp. 248–258\)](#),
- [Nilsson \(1998, chpt. 18.3\)](#)

52. Representing events and time: Situation calculus, event calculus, fluent calculus (including solving the frame problem):

- [Russell & Norvig \(2003, pp. 328–341\)](#),
- [Poole, Mackworth & Goebel \(1998, pp. 281–298\)](#),
- [Nilsson \(1998, chpt. 18.2\)](#)

53. Causal calculus:

- [Poole, Mackworth & Goebel \(1998, pp. 335–337\)](#)

54. Representing knowledge about knowledge: Belief calculus, modal logics:

- [Russell & Norvig \(2003, pp. 341–344\)](#),
- [Poole, Mackworth & Goebel \(1998, pp. 275–277\)](#)

55. Default reasoning, Frame problem logic, non-monotonic logics, circumscription, closed world assumption, abduction:

- [Russell & Norvig \(2003\)](#), pp. 354–360)
- [Poole, Mackworth & Goebel \(1998\)](#), pp. 248–256, 323–335)
- [Luger & Stubblefield \(2004\)](#), pp. 335–363)
- [Nilsson \(1998\)](#), ~18.3.3)

(Poole *et al.* places abduction under "default reasoning". Luger *et al.* places this under "uncertain reasoning").

56. Breadth of commonsense knowledge:

- [Russell & Norvig \(2003\)](#), p. 21),
- [Crevier \(1993\)](#), pp. 113–114),
- [Moravec \(1988\)](#), p. 13),
- [Lenat & Guha \(1989\)](#), Introduction)

57. Smoliar & Zhang (1994).

58. Neumann & Möller (2008).

59. Kuperman, Reichley & Bailey (2006).

60. McGarry (2005).

61. Bertini, Del Bimbo & Torniai (2006).

62. Planning:

- [Russell & Norvig \(2003\)](#), pp. 375–459)
- [Poole, Mackworth & Goebel \(1998\)](#), pp. 281–316)
- [Luger & Stubblefield \(2004\)](#), pp. 314–329)
- [Nilsson \(1998\)](#), chpt. 10.1–2, 22)

Information value theory:

- [Russell & Norvig \(2003\)](#), pp. 600–604)

63. Classical planning:

- [Russell & Norvig \(2003\)](#), pp. 375–430)
- [Poole, Mackworth & Goebel \(1998\)](#), pp. 281–315)
- [Luger & Stubblefield \(2004\)](#), pp. 314–329)
- [Nilsson \(1998\)](#), chpt. 10.1–2, 22)

64. Planning and acting in non-deterministic domains: conditional planning, execution monitoring, replanning and continuous planning:

- [Russell & Norvig \(2003\)](#), pp. 430–449)

gent planning and emergent behavior:

- [Russell & Norvig \(2003\)](#), pp. 449–455)

66. Turing (1950).

67. Solomonoff (1956).

68. Russell & Norvig (2003), pp. 649–788.

69. Learning:

- [Russell & Norvig \(2003\)](#), pp. 649–788)
- [Poole, Mackworth & Goebel \(1998\)](#), pp. 397–438)
- [Luger & Stubblefield \(2004\)](#), pp. 385–542)
- [Nilsson \(1998\)](#), chpt. 3.3, 10.3, 17.5, 20)

70. Reinforcement learning:

- [Russell & Norvig \(2003\)](#), pp. 763–788)
- [Luger & Stubblefield \(2004\)](#), pp. 442–449)

71. The Economist (2016).

72. Jordan & Mitchell (2015).

73. Natural language processing (NLP):

- [Russell & Norvig \(2003\)](#), pp. 790–831)
- [Poole, Mackworth & Goebel \(1998\)](#), pp. 91–104)
- [Luger & Stubblefield \(2004\)](#), pp. 591–632)

74. Applications of NLP:

- [Russell & Norvig \(2003\)](#), pp. 840–857)
- [Luger & Stubblefield \(2004\)](#), pp. 623–630)

75. Modern statistical approaches to NLP:

- [Cambria & White \(2014\)](#)

76. Vincent (2019).

77. Machine perception:

- [Russell & Norvig \(2003\)](#), pp. 537–581, 863–898)
- [Nilsson \(1998\)](#), ~chpt. 6)

78. Speech recognition:

- [Russell & Norvig \(2003\)](#), pp. 568–578)

79. Object recognition:

- [Russell & Norvig \(2003\)](#), pp. 885–892)

80. Computer vision:

- [Russell & Norvig \(2003\)](#), pp. 863–898)
- [Nilsson \(1998\)](#), chpt. 6)

81. Robotics:

- [Russell & Norvig \(2003, pp. 90, 193–219\)](#)
- [Poole, Mackworth & Goebel \(1998, pp. 443–460\)](#)

82. Robotic mapping and Localization:

- [Russell & Norvig \(2003, pp. 908–915\)](#)
- [Cadena et al. \(2016\)](#)

83. Motion planning and configuration space:

- [Russell & Norvig \(2003, pp. 916–932\)](#)
- [Tecuci \(2012\)](#)

84. MIT AIL (2014).

85. Affective computing:

- [Thro \(1993\)](#)
- [Edelson \(1991\)](#)
- [Tao & Tan \(2005\)](#)
- [Scassellati \(2002\)](#)

86. Waddell (2018).

87. Poria et al. (2017).

88. *The Society of Mind*:

- [Minsky \(1986\)](#)

Moravec's "golden spike":

- [Moravec \(1988, p. 20\)](#)

Multi-agent systems, hybrid intelligent systems, agent architectures, cognitive architecture:

- [Russell & Norvig \(2003, pp. 27, 932, 970–972\)](#)
- [Nilsson \(1998, chpt. 25\)](#)

89. Domingos (2015), Chpt. 9.

90. Artificial brain as an approach to AGI:

- [Russell & Norvig \(2003, p. 957\)](#)
- [Crevier \(1993, pp. 271 & 279\)](#)
- [Goertzel et al. \(2010\)](#)

A few of the people who make some form of the argument:

- [Moravec \(1988, p. 20\)](#)
- [Kurzweil \(2005, p. 262\)](#)
- [Hawkins & Blakeslee \(2005\)](#)

algorithms:

- [Russell & Norvig \(2003, pp. 59–189\)](#)
- [Poole, Mackworth & Goebel \(1998, pp. 113–163\)](#)
- [Luger & Stubblefield \(2004, pp. 79–164, 193–219\)](#)
- [Nilsson \(1998, chpt. 7–12\)](#)

92. Forward chaining, backward chaining, Horn clauses, and logical deduction as search:

- [Russell & Norvig \(2003, pp. 217–225, 280–294\)](#)
- [Poole, Mackworth & Goebel \(1998, pp. ~46–52\)](#)
- [Luger & Stubblefield \(2004, pp. 62–73\)](#)
- [Nilsson \(1998, chpt. 4.2, 7.2\)](#)

93. State space search and planning:

- [Russell & Norvig \(2003, pp. 382–387\)](#)
- [Poole, Mackworth & Goebel \(1998, pp. 298–305\)](#)
- [Nilsson \(1998, chpt. 10.1–2\)](#)

94. Moving and configuration space:

- [Russell & Norvig \(2003, pp. 916–932\)](#)

95. Uninformed searches (breadth first search, depth first search and general state space search):

- [Russell & Norvig \(2003, pp. 59–93\)](#)
- [Poole, Mackworth & Goebel \(1998, pp. 113–132\)](#)
- [Luger & Stubblefield \(2004, pp. 79–121\)](#)
- [Nilsson \(1998, chpt. 8\)](#)

96. Heuristic or informed searches (e.g., greedy best first and A^{*}):

- [Russell & Norvig \(2003, pp. 94–109\)](#)
- [Poole, Mackworth & Goebel \(1998, pp. 132–147\)](#)
- [Poole & Mackworth \(2017, Section 3.6\)](#)
- [Luger & Stubblefield \(2004, pp. 133–150\)](#)

97. Tecuci (2012).

98. Optimization searches:

- Russell & Norvig (2003, pp. 110–120, 120–129)
- Poole, Mackworth & Goebel (1998, pp. 56–163)
- Luger & Stubblefield (2004, pp. 127–133)

99. Genetic programming and genetic algorithms:

- Luger & Stubblefield (2004, pp. 509–530)
- Nilsson (1998, chpt. 4.2)

100. Artificial life and society based learning:

- Luger & Stubblefield (2004, pp. 530–541)
- Merkle & Middendorf (2013)

101. Logic:

- Russell & Norvig (2003, pp. 194–310)
- Luger & Stubblefield (2004, pp. 35–77)
- Nilsson (1998, chpt. 13–16)

102. Satplan:

- Russell & Norvig (2003, pp. 402–407)
- Poole, Mackworth & Goebel (1998, pp. 300–301)
- Nilsson (1998, chpt. 21)

103. Explanation based learning, relevance based learning, inductive logic programming, case based reasoning:

- Russell & Norvig (2003, pp. 678–710)
- Poole, Mackworth & Goebel (1998, pp. 414–416)
- Luger & Stubblefield (2004, pp. ~422–442)
- Nilsson (1998, chpt. 10.3, 17.5)

104. Propositional logic:

- Russell & Norvig (2003, pp. 204–233)
- Luger & Stubblefield (2004, pp. 45–50)
- Nilsson (1998, chpt. 13)

under logic and features such as



- Russell & Norvig (2003, pp. 240–310)
- Poole, Mackworth & Goebel (1998, pp. 268–275)
- Luger & Stubblefield (2004, pp. 50–62)
- Nilsson (1998, chpt. 15)

106. Fuzzy logic:

- Russell & Norvig (2003, pp. 526–527)
- Scientific American (1999)

107. Abe, Jair Minoro; Nakamatsu, Kazumi (2009). "Multi-agent Systems and Paraconsistent Knowledge". *Knowledge Processing and Decision Making in Agent-Based Systems* (https://link.springer.com/chapter/10.1007/978-3-540-88049-3_5). *Studies in Computational Intelligence*. Vol. 170. Springer Berlin Heidelberg. pp. 101–121. doi:[10.1007/978-3-540-88049-3_5](https://doi.org/10.1007/978-3-540-88049-3_5) (https://doi.org/10.1007%2F978-3-540-88049-3_5). eISSN 1860-9503 (<https://www.worldcat.org/issn/1860-9503>). ISBN 978-3-540-88048-6. ISSN 1860-949X (<https://www.worldcat.org/issn/1860-949X>). Retrieved 2 August 2022.

108. Stochastic methods for uncertain reasoning:

- Russell & Norvig (2003, pp. 462–644)
- Poole, Mackworth & Goebel (1998, pp. 345–395)
- Luger & Stubblefield (2004, pp. 165–191, 333–381)
- Nilsson (1998, chpt. 19)

109. Bayesian networks:

- Russell & Norvig (2003, pp. 492–523)
- Poole, Mackworth & Goebel (1998, pp. 361–381)
- Luger & Stubblefield (2004, pp. ~182–190, ~363–379)
- Nilsson (1998, chpt. 19.3–4)

110. Domingos (2015), chapter 6.

111. Bayesian inference algorithm:

- Russell & Norvig (2003, pp. 504–519)
- Poole, Mackworth & Goebel (1998, pp. 361–381)
- Luger & Stubblefield (2004, pp. ~363–379)
- Nilsson (1998, chpt. 19.4 & 7)

112. [Domingos \(2015\)](#), p. 210.

113. [Bayesian learning](#) and the [expectation–maximization algorithm](#):

- [Russell & Norvig \(2003\)](#), pp. 712–724),
- [Poole, Mackworth & Goebel \(1998\)](#), pp. 424–433),
- [Nilsson \(1998\)](#), chpt. 20)
- [Domingos \(2015\)](#), p. 210)

114. [Bayesian decision theory](#) and [Bayesian decision networks](#):

- [Russell & Norvig \(2003\)](#), pp. 597–600)

115. Stochastic temporal models:

- [Russell & Norvig \(2003\)](#), pp. 537–581)

[Dynamic Bayesian networks](#):

- [Russell & Norvig \(2003\)](#), pp. 551–557)

[Hidden Markov model](#):

- [\(Russell & Norvig 2003\)](#), pp. 549–551)

[Kalman filters](#):

- [Russell & Norvig \(2003\)](#), pp. 551–557)

116. [decision theory](#) and [decision analysis](#):

- [Russell & Norvig \(2003\)](#), pp. 584–597),
- [Poole, Mackworth & Goebel \(1998\)](#), pp. 381–394)

117. [Information value theory](#):

- [Russell & Norvig \(2003\)](#), pp. 600–604)

118. [Markov decision processes](#) and [dynamic decision networks](#):

- [Russell & Norvig \(2003\)](#), pp. 613–631)

119. [Game theory](#) and [mechanism design](#):

- [Russell & Norvig \(2003\)](#), pp. 631–643)

120. Statistical learning methods and [classifiers](#):

- [Russell & Norvig \(2003\)](#), pp. 712–754),
- [Luger & Stubblefield \(2004\)](#), pp. 453–541)

121. [Decision tree](#):

- [Domingos \(2015\)](#), p. 88)
- [Russell & Norvig \(2003\)](#), pp. 653–664),
- [Poole, Mackworth & Goebel \(1998\)](#), pp. 403–408),
- [Luger & Stubblefield \(2004\)](#), pp. 408–417)

[Nearest neighbor algorithm](#):

- [Domingos \(2015\)](#), p. 187)

- [Russell & Norvig \(2003\)](#), pp. 733–736)

123. [kernel methods](#) such as the [support vector machine](#):

- [Domingos \(2015\)](#), p. 88)

- [Russell & Norvig \(2003\)](#), pp. 749–752)

[Gaussian mixture model](#):

- [Russell & Norvig \(2003\)](#), pp. 725–727)

124. [Domingos \(2015\)](#), p. 152.

125. [Naive Bayes classifier](#):

- [Domingos \(2015\)](#), p. 152)

- [Russell & Norvig \(2003\)](#), p. 718)

126. Neural networks:

- [Russell & Norvig \(2003\)](#), pp. 736–748),

- [Poole, Mackworth & Goebel \(1998\)](#), pp. 408–414),

- [Luger & Stubblefield \(2004\)](#), pp. 453–505),

- [Nilsson \(1998\)](#), chpt. 3)

- [Domingos \(2015\)](#), Chapter 4)

127. [Classifier performance](#):

- [van der Walt & Bernard \(2006\)](#)

- [Russell & Norvig \(2009\)](#), 18.12: Learning from Examples: Summary)

128. [Backpropagation](#):

- [Russell & Norvig \(2003\)](#), pp. 744–748),

- [Luger & Stubblefield \(2004\)](#), pp. 467–474),

- [Nilsson \(1998\)](#), chpt. 3.3)

Paul Werbos' introduction of backpropagation to AI:

- [Werbos \(1974\)](#); [Werbos \(1982\)](#)

Automatic differentiation, an essential precursor:

- [Linnainmaa \(1970\)](#); [Griewank \(2012\)](#)

129. [Competitive learning](#), [Hebbian coincidence learning](#), [Hopfield networks](#) and attractor networks:

- [Luger & Stubblefield \(2004\)](#), pp. 474–505)

130. Feedforward neural networks, perceptrons and radial basis networks:
- Russell & Norvig (2003, pp. 739–748, 758)
 - Luger & Stubblefield (2004, pp. 458–467)
131. Schulz & Behnke (2012).
132. Deep learning:
- Goodfellow, Bengio & Courville (2016)
 - Hinton *et al.* (2016)
 - Schmidhuber (2015)
133. Deng & Yu (2014), pp. 199–200.
134. Ciresan, Meier & Schmidhuber (2012).
135. Habibi (2017).
136. Fukushima (2007).
137. Recurrent neural networks, Hopfield nets:
- Russell & Norvig (2003, p. 758)
 - Luger & Stubblefield (2004, pp. 474–505)
 - Schmidhuber (2015)
138. Schmidhuber (2015).
139. Werbos (1988); Robinson & Fallside (1987); Williams & Zipser (1994)
140. Goodfellow, Bengio & Courville (2016); Hochreiter (1991)
141. Hochreiter & Schmidhuber (1997); Gers, Schraudolph & Schraudolph (2002)
142. Russell & Norvig (2009), p. 1.
143. European Commission (2020), p. 1.
144. CNN (2006).
145. Targeted advertising:
- Russell & Norvig (2009, p. 1)
 - Economist (2016)
 - Lohr (2016)
146. Lohr (2016).
147. Smith (2016).
148. Rowinski (2013).
149. Frangoul (2019).
150. Brown (2019).
151. McCorduck (2004), pp. 480–483.
152. Markoff (2011).
153. Google (2016); BBC (2016)
154. Solly (2019).
155. Bowling *et al.* (2015).
156. Sample (2017).
- tis (2020).
2020).
159. Aletras *et al.* (2016).
160. "Intellectual Property and Frontier Technologies" (https://www.wipo.int/about-ip/en/frontier_technologies/). WIPO.
161. "WIPO Technology Trends 2019 – Artificial Intelligence" (https://www.wipo.int/edocs/patentdocs/en/wipo_pub_1055.pdf) (PDF). WIPO. 2019.
162. Turing (1950), p. 1.
163. Turing (1948).
164. Turing's original publication of the Turing test in "Computing machinery and intelligence":
- Turing (1950)
- Historical influence and philosophical implications:
- Haugeland (1985, pp. 6–9)
 - Crevier (1993, p. 24)
 - McCorduck (2004, pp. 70–71)
 - Russell & Norvig (2021, pp. 2 and 984)
165. Russell & Norvig (2021), p. 2–3.
166. Russell & Norvig (2021), p. 3.
167. Maker (2006).
168. McCorduck (2004), pp. 100–101.
169. The intelligent agent paradigm:
- Russell & Norvig (2021, p. 4, chpt. 2)
 - Poole, Mackworth & Goebel (1998, pp. 7–21)
 - Luger & Stubblefield (2004, pp. 235–240)
 - Hutter (2005, pp. 125–126)
- The definition used in this article, in terms of goals, actions, perception and environment, is due to Russell & Norvig (2021, p. 40). Other definitions also include knowledge, learning and autonomy as additional criteria.
170. Russell & Norvig (2021), p. 4.
171. Nilsson (1983), p. 10.
172. Haugeland (1985), pp. 112–117.
173. Physical symbol system hypothesis:
- Newell & Simon (1976, p. 116)
- Historical significance:
- McCorduck (2004, p. 153)
 - Russell & Norvig (2003, p. 18)



174. Moravec's paradox:

- [Moravec \(1988, pp. 15–16\)](#)
- [Minsky \(1986, p. 29\)](#)
- [Pinker \(2007, pp. 190–91\)](#)

175. Dreyfus' critique of AI:

- [Dreyfus \(1972\)](#)
- [Dreyfus & Dreyfus \(1986\)](#)

Historical significance and philosophical implications:

- [Crevier \(1993, pp. 120–132\)](#)
- [McCorduck \(2004, pp. 211–239\)](#)
- [Russell & Norvig \(2003, pp. 950–952\)](#)
- [Fearn \(2007, Chpt. 3\)](#)

176. [Crevier \(1993\)](#), p. 125.

177. [Langley \(2011\)](#).

178. [Katz \(2012\)](#).

179. [Neats vs. scruffies](#), the historic debate:

- [McCorduck \(2004, pp. 421–424, 486–489\)](#)
- [Crevier \(1993, p. 168\)](#)
- [Nilsson \(1983, pp. 10–11\)](#)

A classic example of the "scruffy" approach to intelligence:

- [Minsky \(1986\)](#)

A modern example of neat AI and its aspirations:

- [Domingos \(2015\)](#)

180. [Russell & Norvig \(2003\)](#), pp. 25–26.

181. [Pennachin & Goertzel \(2007\)](#).

182. [Roberts \(2016\)](#).

183. [Russell & Norvig \(2003\)](#), p. 947.

184. [Chalmers \(1995\)](#).

185. [Dennett \(1991\)](#).

186. [Horst \(2005\)](#).

187. [Searle \(1999\)](#).

188. [Searle \(1980\)](#), p. 1.

189. Searle's [Chinese room](#) argument:

- [Searle \(1980\)](#). Searle's original presentation of the thought experiment.
- [Searle \(1999\)](#).

Discussion:

- [Russell & Norvig \(2003, pp. 958–960\)](#)
- [McCorduck \(2004, pp. 443–445\)](#)
- [Crevier \(1993, pp. 269–271\)](#)

rights:

- [. Russell & Norvig \(2003, p. 964\)](#)
- [BBC \(2006\)](#)
- [Maschafilm \(2010\) \(the film Plug & Pray\)](#)

191. [Evans \(2015\)](#).

192. [McCorduck \(2004\)](#), pp. 19–25.

193. [Henderson \(2007\)](#).

194. [Omohundro \(2008\)](#).

195. [Vinge \(1993\)](#).

196. [Russell & Norvig \(2003\)](#), p. 963.

197. [Transhumanism](#):

- [Moravec \(1988\)](#)
- [Kurzweil \(2005\)](#)
- [Russell & Norvig \(2003, p. 963\)](#)

198. [AI as evolution](#):

- [Edward Fredkin](#) is quoted in [McCorduck \(2004, p. 401\)](#)
- [Butler \(1863\)](#)
- [Dyson \(1998\)](#)

199. [Ford & Colvin \(2015\)](#); [McGaughey \(2018\)](#)

200. [IGM Chicago \(2017\)](#).

201. [Arntz, Gregory & Zierahn \(2016\)](#), p. 33.

202. [Lohr \(2017\)](#); [Frey & Osborne \(2017\)](#); [Arntz, Gregory & Zierahn \(2016, p. 33\)](#)

203. [Morgenstern \(2015\)](#).

204. [Mahdawi \(2017\)](#); [Thompson \(2014\)](#)

205. [Harari \(2018\)](#).

206. [Weaponized AI](#):

- [Robitzski \(2018\)](#)
- [Sainato \(2015\)](#)

207. Urbina, Fabio; Lentzos, Filippa; Invernizzi, Cédric; Ekins, Sean (7 March 2022). "[Dual use of artificial-intelligence-powered drug discovery](#)" (<https://www.nature.com/articles/s42256-022-00465-9>). *Nature Machine Intelligence*. 4 (3): 189–191. doi:[10.1038/s42256-022-00465-9](https://doi.org/10.1038/s42256-022-00465-9) (<https://doi.org/10.1038%2Fs42256-022-00465-9>). S2CID 247302391 (<https://api.semanticscholar.org/CorpusID:247302391>). Retrieved 15 March 2022.

208. [CNA \(2019\)](#).

209. [Goffrey \(2008\)](#), p. 17.

210. [Lipartito \(2011, p. 36\)](#); [Goodman & Flaxman \(2017, p. 6\)](#)

211. [Larson & Angwin \(2016\)](#).

212. Dockrill, Peter, *Robots With Flawed Sexist And Racist Decisions, Experimental Shows* (<https://www.sciencealert.com/robots-with-flawed-ai-make-sexist-racist-and-toxic-decisions-experiment-shows>), Science Alert, 27 June 2022
213. Cellan-Jones (2014).
214. Bostrom (2014); Müller & Bostrom (2014); Bostrom (2015)
215. Rubin (2003).
216. Müller & Bostrom (2014).
217. Leaders' concerns about the existential risks of AI:
- Rawlinson (2015)
 - Holley (2015)
 - Gibbs (2014)
 - Churm (2019)
 - Sainato (2015)
218. Funding to mitigate risks of AI:
- Post (2015)
 - Del Prado (2015)
 - Clark (2015a)
 - FastCompany (2015)
219. Leaders who argue the benefits of AI outweigh the risks:
- Thibodeau (2019)
 - Bhardwaj (2018)
220. Arguments that AI is not an imminent risk:
- Brooks (2014)
 - Geist (2015)
 - Madrigal (2015)
 - Lee (2014)
221. Brooks (2014).
222. "Artificial intelligence and copyright" (https://www.wipo.int/wipo_magazine/en/2017/05/article_0003.html). *www.wipo.int*. Retrieved 27 May 2022.
223. Yudkowsky (2008).
224. Anderson & Anderson (2011).
225. AAAI (2014).
226. Wallach (2010).
227. Russell (2019), p. 173.
- tion of AI to mitigate risks:
- Berryhill et al. (2019)
 - Barfield & Pagallo (2018)
 - Iphofen & Kritikos (2019)
 - Wirtz, Weyerer & Geyer (2018)
 - Buiten (2019)
229. Law Library of Congress (U.S.). Global Legal Research Directorate (2019).
230. Kissinger, Henry (1 November 2021). "The Challenge of Being Human in the Age of AI" (<https://www.wsj.com/articles/being-human-artificial-intelligence-ai-chess-antibiotic-philosophy-ethics-bill-of-rights-11635795271>). *The Wall Street Journal*.
231. Buttazzo (2001).
232. Anderson (2008).
233. McCauley (2007).
234. Galvan (1997).
235. "A Visual Guide to the World's Military Budgets" (<https://www.bloomberg.com/news/features/2022-03-11/the-largest-militaries-visualized>). *Bloomberg.com*. 11 March 2022. Retrieved 11 May 2022.
236. Kharpal, Arjun (21 July 2017). "China wants to be a \$150 billion world leader in AI in less than 15 years" (<https://www.cnbc.com/2017/07/21/china-ai-world-leader-by-2030.html>). *CNBC*. Retrieved 11 May 2022.
237. Radina Gigova (2 September 2017). "Who Putin thinks will rule the world" (<https://www.cnn.com/2017/09/01/world/putin-artificial-intelligence-will-rule-world/index.html>). *CNN*. Retrieved 18 May 2022.
238. "In Ukraine, A.I. is going to war" (<https://fortune.com/2022/03/01/russia-ukraine-invasion-war-a-i-artificial-intelligence/>). *Fortune*. Retrieved 11 May 2022.
239. "AI Is Already Learning from Russia's War in Ukraine, DOD Says" (<https://www.defenseone.com/technology/2022/04/ai-already-learning-russias-war-ukraine-dod-says/365978/>). *Defense One*. Retrieved 11 May 2022.
240. "A.I. drones used in the Ukraine war raise fears of killer robots wreaking havoc across future battlefields" (<https://fortune.com/2022/03/29/artificial-intelligence-drones-autonomous-weapons-loitering-munitions-slaughterbots-ukraine-war/>). *Fortune*. Retrieved 11 May 2022.

241. Vincent, James (21 August 2017). "Musk and AI leaders call for a ban on killer robots" (<https://www.theverge.com/2017/8/2/16177828/killer-robots-ban-elon-musk-unpetition>). *The Verge*. Retrieved 11 May 2022.
242. Kerry, Cameron F. (10 February 2020). "Protecting privacy in an AI-driven world" (<https://www.brookings.edu/research/protecting-privacy-in-an-ai-driven-world/>). *Brookings*. Retrieved 11 May 2022.
243. "CYBERSECURITY | CISA" (<https://www.cisa.gov/cybersecurity>). www.cisa.gov. Retrieved 11 May 2022.
244. Inkster, Nigel (4 November 2021), Cornish, Paul (ed.), "Semi-Formal Diplomacy" (<https://dx.doi.org/10.1093/oxfordhb/9780198800682.013.33>), *The Oxford Handbook of Cyber Security*, Oxford University Press, pp. 530–542, doi:10.1093/oxfordhb/9780198800682.013.33 (<https://doi.org/10.1093%2Foxfordhb%2F9780198800682.013.33>), ISBN 978-0-19-880068-2, retrieved 11 May 2022
245. "Connect the Dots on State-Sponsored Cyber Incidents – Compromise of the Czech foreign minister's computer" (<https://www.cfr.org/index.php/cyber-operations/compromise-czech-foreign-ministers-computer>). *Council on Foreign Relations*. Retrieved 11 May 2022.
246. Kadlecová, Lucie; Meyer, Nadia; Cos, Rafaël; Ravinet, Pauline (2020). "Cyber Security: Mapping the Role of Science Diplomacy in the Cyber Field". In: Young, M., T. Flink, e. Dall. *Science Diplomacy in the Making: Case-based insights from the S4D4C project*.
247. Canada, Global Affairs (19 November 2019). "European innovation partnerships – A guide for Canadian small and medium enterprises" (<https://www.tradecommissioner.gc.ca/guides/eu-innovation-ue.aspx?lang=eng>). GAC. Retrieved 11 May 2022.
248. "Homepage" (https://www.bsi.bund.de/EN/Home/home_node.html;jsessionid=468D98DD1E7E6C20643D9C468F8808F.internet461). *Federal Office for Information Security*. Retrieved 11 May 2022.
249. "Die Künstliche Intelligenz der Bundesregierung" (<https://www.bundesregierung.de/breg-de/service/publikationen/strategische-kuenstliche-intelligenz-der-bundesregierung-2018-1551264>). *Federal Ministry for Economic Affairs and Energy*: 6. 2018 – via Economy and climate protection (BMWK).
250. "Europe counts on Luxembourg's expertise" (<https://www.san-francisco.investinluxembourg.us/news/europe-counts-on-luxembourgs-expertise/>). *San Francisco*. Retrieved 11 May 2022.
251. "Cyber Security Europe | Cyber security insight for boardroom and c-suite executives" (<https://www.cseurope.info/>). *Cyber Security Europe*. Retrieved 11 May 2022.
252. EMK, SU. "ECHO Network" (<https://echonet-work.eu/>). Retrieved 11 May 2022.
253. "SPARTA Consortium" (<https://www.cybersecurityintelligence.com/sparta-consortium-5594.html>). www.cybersecurityintelligence.com. Retrieved 11 May 2022.
254. "Shields Up | CISA" (<https://www.cisa.gov/shields-up>). www.cisa.gov. Retrieved 11 May 2022.
255. Tegler, Eric. "The Vulnerability of AI Systems May Explain Why Russia Isn't Using Them Extensively in Ukraine" (<https://www.forbes.com/sites/erictegler/2022/03/16/the-vulnerability-of-artificial-intelligence-systems-may-explain-why-they-havent-been-used-extensively-in-ukraine/>). *Forbes*. Retrieved 11 May 2022.
256. "CyberAI Project" (<https://cset.georgetown.edu/cyberai-project/>). *Center for Security and Emerging Technology*. Retrieved 11 May 2022.
257. "AI Weekly: The Russia-Ukraine conflict is a test case for AI in warfare" (<https://venturebeat.com/2022/03/04/ai-weekly-the-russia-ukraine-conflict-is-a-test-case-for-ai-in-warfare/>). *VentureBeat*. 4 March 2022. Retrieved 11 May 2022.
258. Kamarck, Elaine (29 November 2018). "Malevolent soft power, AI, and the threat to democracy" (<https://www.brookings.edu/research/malevolent-soft-power-ai-and-the-threat-to-democracy/>). *Brookings*. Retrieved 18 May 2022.

259. Guglielmi, Giorgia (28 October 2020). "next-generation bots interfering with election" (<https://www.nature.com/articles/a41586-020-03034-5>). *Nature*. **587** (7832): 21. Bibcode:2020Natur.587...21G (<https://ui.adsabs.harvard.edu/abs/2020Natur.587...21G>). doi:10.1038/d41586-020-03034-5 (<https://doi.org/10.1038/d41586-020-03034-5>). PMID 33116324 (<https://pubmed.ncbi.nlm.nih.gov/33116324>). S2CID 226052075 (<https://api.semanticscholar.org/CorpusID:226052075>).
260. HosseiniMardi, Homa; Ghasemian, Amir; Clauset, Aaron; Mobius, Markus; Rothschild, David M.; Watts, Duncan J. (10 August 2021). "Examining the consumption of radical content on YouTube" (<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC8364190>). *Proceedings of the National Academy of Sciences*. **118** (32): e2101967118. doi:10.1073/pnas.2101967118 (<https://doi.org/10.1073/pnas.2101967118>). ISSN 0027-8424 (<https://www.worldcat.org/issn/0027-8424>). PMC 8364190 (<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC8364190>). PMID 34341121 (<https://pubmed.ncbi.nlm.nih.gov/34341121>).
261. "Explained: Here is how YouTube recommendations work" (<https://indianexpress.com/article/explained/youtube-recommendations-explained-7587523/>). *The Indian Express*. 26 October 2021. Retrieved 18 May 2022.
262. You Won't Believe What Obama Says In This Video! 😊 (<https://www.youtube.com/watch?v=cQ54GDm1eL0>), retrieved 18 May 2022
263. "Where Are The Deepfakes In This Presidential Election?" (<https://www.npr.org/2020/10/01/918223033/where-are-the-deepfakes-in-this-presidential-election>). *NPR.org*. Retrieved 18 May 2022.
264. PhD, Slava Polonski (4 February 2018). "Artificial intelligence can save democracy, unless it destroys it first" (<https://medium.com/@slavaxyz/artificial-intelligence-can-save-democracy-unless-it-destroys-it-first-7b1257cb4285>). *Medium*. Retrieved 18 May 2022.
265. "Science Diplomacy and Future Worlds" (<https://www.sciediplomacy.org/editorial/2018/science-diplomacy-and-future-worlds>). *Science & Diplomacy*. Retrieved 18 May 2022.
- "AI Revenue Service (IRS)" (<https://dx.doi.org/10.4135/9781412956260.n436>), *Encyclopedia of Business Ethics and Society*, Thousand Oaks, California: SAGE Publications, Inc., 2008, doi:10.4135/9781412956260.n436 (<https://doi.org/10.4135%2F9781412956260.n436>), ISBN 9781412916523, retrieved 19 May 2022
267. Ga-young, Park (9 October 2021). "Korea's edu tech startup Riid acquires Japanese mobile app distributor" (<http://www.koreaherald.com/view.php?ud=20211009000064>). *The Korea Herald*. Retrieved 19 May 2022.
268. "You are being redirected..." (<https://www.analyticsinsight.net/will-artificial-intelligence-replace-teachers-in-the-future/>) www.analyticsinsight.net. Retrieved 19 May 2022.
269. "AI early diagnosis could save heart and cancer patients" (<https://www.bbc.com/news/health-42357257>). *BBC News*. 2 January 2018. Retrieved 19 May 2022.
270. Davenport, Thomas; Kalakota, Ravi (June 2019). "The potential for artificial intelligence in healthcare" (<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC6616181>). *Future Healthcare Journal*. **6** (2): 94–98. doi:10.7861/futurehosp.6-2-94 (<https://doi.org/10.7861/futurehosp.6-2-94>). ISSN 2514-6645 (<https://www.worldcat.org/issn/2514-6645>). PMC 6616181 (<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC6616181>). PMID 31363513 (<https://pubmed.ncbi.nlm.nih.gov/31363513>).
271. Jain J, Pragya. "AI and the Future of Work in the United States" (<https://www.american.edu/sis/centers/security-technology/ai-and-the-future-of-work-in-the-united-states.cfm>). *American University*. Retrieved 19 May 2022.
272. "BASF and its partners publish results for 'Pragati', world's first sustainable castor bean program" (<https://dx.doi.org/10.1016/j.fopow.2022.01.019>). *Focus on Powder Coatings*. **2022** (2): 5. February 2022. doi:10.1016/j.fopow.2022.01.019 (<https://doi.org/10.1016/j.fopow.2022.01.019>). ISSN 1364-5439 (<https://www.worldcat.org/issn/1364-5439>). S2CID 246561954 (<https://api.semanticscholar.org/CorpusID:246561954>).

273. "AI in 2020: From Experimentation to Adoption" (<https://www.ibm.com/blogs/think/2020/01/ai-in-2020-from-experimentation-to-adoption/>). *THINK Blog*. 3 January 2020. Retrieved 19 May 2022.
274. "Artificial Intelligence Market Size, Share, Trends, Opportunities & Forecast" (<https://www.verifiedmarketresearch.com/product/global-artificial-intelligence-market-size-and-forecast-to-2025/>). *Verified Market Research*. Retrieved 19 May 2022.
275. "European Commission.: Ethics guidelines for trustworthy AI. EC HLEG" (<https://digital-strategy.ec.europa.eu/en/library/ethics-guidelines-trustworthy-ai>). 2019.
276. Curtis, Caitlin; Gillespie, Nicole; Lockey, Steven (24 May 2022). "AI-deploying organizations are key to addressing 'perfect storm' of AI risks" (<https://doi.org/10.1007/s43681-022-00163-7>). *AI and Ethics*: 1–9. doi:10.1007/s43681-022-00163-7 (<https://doi.org/10.1007%2Fs43681-022-00163-7>). ISSN 2730-5961 (<https://www.worldcat.org/issn/2730-5961>). PMC 9127285 (<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC9127285>). PMID 35634256 (<https://pubmed.ncbi.nlm.nih.gov/35634256>).
277. "Fakers Who Realize That They're the Real Thing" (<https://dx.doi.org/10.4324/9780203508305-31>), *The New York Times Television Reviews 2000*, Routledge, pp. 423–445, 5 June 2003, doi:10.4324/9780203508305-31 (<https://doi.org/10.4324%2F9780203508305-31>), ISBN 9780203508305, retrieved 19 May 2022
278. Barattero, Alberto (2018). "The People Vs Tech: How the Internet is Killing Democracy (and How We Save It)" (<https://dx.doi.org/10.5840/tiw20185221>). *The Incarnate Word*. 5 (2): 204–207. doi:10.5840/tiw20185221 (<https://doi.org/10.5840%2Ftiw20185221>). ISSN 2150-9824 (<https://www.worldcat.org/issn/2150-9824>).
279. Tapsell, Paul (December 2021), "He Tohu: Tipping Point" (https://dx.doi.org/10.7810/9781988587585_1), *Kāinga: People, Land, Belonging*, Bridget Williams Books, pp. 11–24, doi:10.7810/9781988587585_1 (https://doi.org/10.7810%2F9781988587585_1), ISBN 9781988587585, S2CID 245941972 (<https://api.semanticscholar.org/CorpusID:245941972>), retrieved 19 May 2022
280. 'gorithms make safe decisions in affairs? – Diplo' (<https://www.diplomacy.edu/blog/will-algorithms-make-safe-decisions-foreign-affairs/>). 17 December 2019. Retrieved 19 May 2022.
281. "How AI Is Running China's Foreign Policy" (<https://analyticsindiamag.com/how-ai-is-running-chinas-foreign-policy/>). *Analytics India Magazine*. 31 July 2018. Retrieved 19 May 2022.
282. "United Nations Activities on Artificial Intelligence (AI)" (<https://www.itu.int:443/en/publications/gs/Pages/publications.aspx>). *ITU*. Retrieved 19 May 2022.
283. "Engineering Diplomacy: How AI and Human Augmentation Could Remake the Art of Foreign Relations" (<https://www.scienceanddiplomacy.org/perspective/2022/engineering-diplomacy-how-ai-and-human-augmentation-could-remake-art-foreign>). *Science & Diplomacy*. Retrieved 19 May 2022.
284. Aylett-Bullock, Joseph; Cuesta-Lazaro, Carolina; Quera-Bofarull, Arnau; Katta, Anjali; Pham, Katherine Hoffmann; Hoover, Benjamin; Strobelt, Hendrik; Jimenez, Rebeca Moreno; Sedgewick, Aidan; Evers, Egmond Samir; Kennedy, David (6 July 2021). "Operational response simulation tool for epidemics within refugee and IDP settlements" (<https://www.medrxiv.org/content/10.1101/2021.01.27.21250611v2>): 2021.01.27.21250611. doi:10.1101/2021.01.27.21250611 (<https://doi.org/10.1101%2F2021.01.27.21250611>). S2CID 231722795 (<https://api.semanticscholar.org/CorpusID:231722795>).
285. Witmer, Frank D. W. (3 May 2015). "Remote sensing of violent conflict: eyes from above" (<https://doi.org/10.1080/01431161.2015.1035412>). *International Journal of Remote Sensing*. 36 (9): 2326–2352. Bibcode:2015IJRS...36.2326W (<https://ui.adsabs.harvard.edu/abs/2015IJRS...36.2326W>). doi:10.1080/01431161.2015.1035412 (<https://doi.org/10.1080%2F01431161.2015.1035412>). ISSN 0143-1161 (<https://www.worldcat.org/issn/0143-1161>). S2CID 140656194 (<https://api.semanticscholar.org/CorpusID:140656194>).

286. Wu, Jeff; Ouyang, Long; Ziegler, D
Stiennon, Nisan; Lowe, Ryan; Lei
Christiano, Paul (27 September 2021). "Recursively Summarizing Books with Human Feedback". arXiv:2109.10862 ([http://arxiv.org/abs/2109.10862](https://arxiv.org/abs/2109.10862)) [cs.CL (<https://arxiv.org/archive/cs.CL>)].
287. "What the American Interpreter Got Wrong at Tense US-China Talks in Alaska" (<https://www.vice.com/en/article/epdy4p/us-china-interpreter-mistranslation>). www.vice.com. Retrieved 19 May 2022.

'a, Ivano; Lavelli, Alberto; Aiolfi, Fabio
January 2022). "An introduction to Deep Learning in Natural Language Processing: Models, techniques, and tools" (<https://www.sciencedirect.com/science/article/pii/S0925231221010997>). *Neurocomputing*. 470: 443–456. doi:10.1016/j.neucom.2021.05.103 (<https://doi.org/10.1016%2Fj.neucom.2021.05.103>). ISSN 0925-2312 (<https://www.worldcat.org/issn/0925-2312>). S2CID 238835461 (<https://api.semanticscholar.org/CorpusID:238835461>).

References

Attribution

- This article incorporates text from a free content work. Licensed under C-BY-SA 3.0 IGO. Text taken from *UNESCO Science Report: the Race Against Time for Smarter Development*. (<https://unesdoc.unesco.org/ark:/48223/pf0000377433/PDF/377433eng.pdf.multi>), Schneegans, S., T. Straza and J. Lewis (eds), UNESCO.

AI textbooks

These were the four the most widely used AI textbooks in 2008:

- Luger, George; Stubblefield, William (2004). *Artificial Intelligence: Structures and Strategies for Complex Problem Solving* (<https://archive.org/details/artificialintell0000luge>) (5th ed.). Benjamin/Cummings. ISBN 978-0-8053-4780-7. Archived (<https://web.archive.org/web/20200726220613/https://archive.org/details/artificialintell0000luge>) from the original on 26 July 2020. Retrieved 17 December 2019.
- Nilsson, Nils (1998). *Artificial Intelligence: A New Synthesis* (<https://archive.org/details/artificialintell0000nils>). Morgan Kaufmann. ISBN 978-1-55860-467-4. Archived (<https://web.archive.org/web/20200726131654/https://archive.org/details/artificialintell0000nils>) from the original on 26 July 2020. Retrieved 18 November 2019.
- Russell, Stuart J.; Norvig, Peter (2003). *Artificial Intelligence: A Modern Approach* (<http://aima.cs.berkeley.edu/>) (2nd ed.), Upper Saddle River, New Jersey: Prentice Hall, ISBN 0-13-790395-2.
- Poole, David; Mackworth, Alan; Goebel, Randy (1998). *Computational Intelligence: A Logical Approach* (<https://archive.org/details/computationalint00pool>). New York: Oxford University Press. ISBN 978-0-19-510270-3. Archived (<https://web.archive.org/web/20200726131436/http://archive.org/details/computationalint00pool>) from the original on 26 July 2020. Retrieved 22 August 2020.

Later editions.

- Russell, Stuart J.; Norvig, Peter (2009). *Artificial Intelligence: A Modern Approach* (3rd ed.). Upper Saddle River, New Jersey: Prentice Hall. ISBN 978-0-13-604259-4..
- Poole, David; Mackworth, Alan (2017). *Artificial Intelligence: Foundations of Computational Agents* (<http://artint.info/index.html>) (2nd ed.). Cambridge University Press. ISBN 978-1-107-19539-4.

The two most widely used textbooks in the field of Computer Science are:

Explorer (<https://opensyllabus.org/resource?field=id&Computer+Science>)

- Russell, Stuart J.; Norvig, Peter (2021). *Artificial Intelligence: A Modern Approach* (4th ed.). Hoboken: Pearson. ISBN 9780134610993. LCCN 20190474 (<https://lccn.loc.gov/20190474>).
- Knight, Kevin; Rich, Elaine (1 January 2010). *Artificial Intelligence* (3rd ed.). Mc Graw Hill India. ISBN 9780070087705.

History of AI

- Crevier, Daniel (1993). *AI: The Tumultuous Search for Artificial Intelligence*. New York, NY: BasicBooks. ISBN 0-465-02997-3..
- McCorduck, Pamela (2004), *Machines Who Think* (2nd ed.), Natick, MA: A. K. Peters, Ltd., ISBN 1-56881-205-1.
- Newquist, HP (1994). *The Brain Makers: Genius, Ego, And Greed In The Quest For Machines That Think*. New York: Macmillan/SAMS. ISBN 978-0-672-30412-5.
- Nilsson, Nils (2009). *The Quest for Artificial Intelligence: A History of Ideas and Achievements*. New York: Cambridge University Press. ISBN 978-0-521-12293-1.

Other sources

- Werbos, P. J. (1988), "Generalization of backpropagation with application to a recurrent gas market model" (<https://zenodo.org/record/1258627>), *Neural Networks*, 1 (4): 339–356, doi:10.1016/0893-6080(88)90007-X (<https://doi.org/10.1016%2F0893-6080%2888%2990007-X>)
- Gers, Felix A.; Schraudolph, Nicol N.; Schraudolph, Jürgen (2002). "Learning Precise Timing with LSTM Recurrent Networks" (<http://www.jmlr.org/papers/volume3/gers02a/gers02a.pdf>) (PDF). *Journal of Machine Learning Research*. 3: 115–143. Retrieved 13 June 2017.
- Deng, L.; Yu, D. (2014). "Deep Learning: Methods and Applications" (<http://research.microsoft.com/pubs/209355/DeepLearning-NowPublishing-Vol7-SIG-039.pdf>) (PDF). *Foundations and Trends in Signal Processing*. 7 (3–4): 1–199. doi:10.1561/2000000039 (<https://doi.org/10.1561%2F2000000039>). Archived (<https://web.archive.org/web/20160314152112/http://research.microsoft.com/pubs/209355/DeepLearning-NowPublishing-Vol7-SIG-039.pdf>) (PDF) from the original on 14 March 2016. Retrieved 18 October 2014.
- Schulz, Hannes; Behnke, Sven (1 November 2012). "Deep Learning" (<https://www.researchgate.net/publication/230690795>). *KI – Künstliche Intelligenz*. 26 (4): 357–363. doi:10.1007/s13218-012-0198-z (<https://doi.org/10.1007%2Fs13218-012-0198-z>). ISSN 1610-1987 (<https://www.worldcat.org/issn/1610-1987>). S2CID 220523562 (<https://api.semanticscholar.org/CorpusID:220523562>).
- Fukushima, K. (2007). "Neocognitron" (<https://doi.org/10.4249%2Fscholarpedia.1717>). *Scholarpedia*. 2 (1): 1717. Bibcode:2007SchpJ...2.1717F (<https://ui.adsabs.harvard.edu/abs/2007SchpJ...2.1717F>). doi:10.4249/scholarpedia.1717 (<https://doi.org/10.4249%2Fscholarpedia.1717>). was introduced by Kunihiko Fukushima in 1980.
- Habibi, Aghdam, Hamed (30 May 2017). *Guide to convolutional neural networks : a practical application to traffic-sign detection and classification*. Heravi, Elnaz Jahani. Cham, Switzerland. ISBN 9783319575490. OCLC 987790957 (<https://www.worldcat.org/oclc/987790957>).
- Ciresan, D.; Meier, U.; Schmidhuber, J. (2012). "Multi-column deep neural networks for image classification". *2012 IEEE Conference on Computer Vision and Pattern Recognition*. pp. 3642–3649. arXiv:1202.2745 (<https://arxiv.org/abs/1202.2745>). doi:10.1109/cvpr.2012.6248110 (<https://doi.org/10.1109%2Fcvpr.2012.6248110>). ISBN 978-1-4673-1228-8. S2CID 2161592 (<https://api.semanticscholar.org/CorpusID:2161592>).

- "From not working to neural networks: The Economist's special report on AI". *The Economist*. 2016. Archived (<https://web.archive.org/web/20161231203934/https://www.economist.com/news/special-report/21700756-artificial-intelligence-boom-based-old-idea-modern-twist-not>) from the original on 31 December 2016. Retrieved 26 April 2018.
- Thompson, Derek (23 January 2014). "What Jobs Will the Robots Take?" (<https://www.theatlantic.com/business/archive/2014/01/what-jobs-will-the-robots-take/283239/>). *The Atlantic*. Archived (<https://web.archive.org/web/20180424202435/https://www.theatlantic.com/business/archive/2014/01/what-jobs-will-the-robots-take/283239/>) from the original on 24 April 2018. Retrieved 24 April 2018.
- Scassellati, Brian (2002). "Theory of mind for a humanoid robot". *Autonomous Robots*. 12 (1): 13–24. doi:[10.1023/A:1013298507114](https://doi.org/10.1023/A:1013298507114) (<https://doi.org/10.1023%2FA%1013298507114>). S2CID 1979315 (<https://api.semanticscholar.org/CorpusID:1979315>).
- Sample, Ian (14 March 2017). "Google's DeepMind makes AI program that can learn like a human" (<https://www.theguardian.com/global/2017/mar/14/googles-deepmind-makes-ai-program-that-can-learn-like-a-human>). *The Guardian*. Archived (<https://web.archive.org/web/20180426212908/https://www.theguardian.com/global/2017/mar/14/googles-deepmind-makes-ai-program-that-can-learn-like-a-human>) from the original on 26 April 2018. Retrieved 26 April 2018.
- Heath, Nick (11 December 2020). "What is AI? Everything you need to know about Artificial Intelligence" (<https://www.zdnet.com/article/what-is-ai-everything-you-need-to-know-about-artificial-intelligence/>). *ZDNet*. Retrieved 1 March 2021.
- Bowling, Michael; Burch, Neil; Johanson, Michael; Tammelin, Oskari (9 January 2015). "Heads-up limit hold'em poker is solved" (<https://www.science.org/doi/10.1126/science.1259433>). *Science*. 347 (6218): 145–149. Bibcode:2015Sci...347..145B (<https://ui.adsabs.harvard.edu/abs/2015Sci...347..145B>). doi:[10.1126/science.1259433](https://doi.org/10.1126/science.1259433) (<https://doi.org/10.1126%2Fscience.1259433>). ISSN 0036-8075 (<https://www.worldcat.org/issn/0036-8075>). PMID 25574016 (<https://pubmed.ncbi.nlm.nih.gov/25574016>). S2CID 3796371 (<https://api.semanticscholar.org/CorpusID:3796371>).
- Solly, Meilan (15 July 2019). "This Poker-Playing A.I. Knows When to Hold 'Em and When to Fold 'Em" (<https://www.smithsonianmag.com/smart-news/poker-playing-ai-knows-when-hold-em-when-fold-em-180972643/>). *Smithsonian*.
- "Artificial intelligence: Google's AlphaGo beats Go master Lee Se-dol" (<https://www.bbc.com/news/technology-35785875>). *BBC News*. 12 March 2016. Archived (<https://web.archive.org/web/20160826103910/http://www.bbc.com/news/technology-35785875>) from the original on 26 August 2016. Retrieved 1 October 2016.
- Rowinski, Dan (15 January 2013). "Virtual Personal Assistants & The Future Of Your Smartphone [Infographic]" (<http://readwrite.com/2013/01/15/virtual-personal-assistants-the-future-of-your-smartphone-infographic>). *ReadWrite*. Archived (<https://web.archive.org/web/20151222083034/http://readwrite.com/2013/01/15/virtual-personal-assistants-the-future-of-your-smartphone-infographic>) from the original on 22 December 2015.
- Manyika, James (2022). "Getting AI Right: Introductory Notes on AI & Society" (<https://www.americanacademyofartsandsciences.org/publication/getting-ai-right-introductory-notes-ai-society>). *Daedalus*. 151 (2): 5–27. doi:[10.1162/daed_e_01897](https://doi.org/10.1162/daed_e_01897) (https://doi.org/10.1162%2Fdaed_e_01897). S2CID 248377878 (<https://api.semanticscholar.org/CorpusID:248377878>). Retrieved 5 May 2022.
- Markoff, John (16 February 2011). "Computer Wins on 'Jeopardy!': Trivial, It's Not" (<https://www.nytimes.com/2011/02/17/science/17jeopardy-watson.html>). *The New York Times*. Archived (<https://web.archive.org/web/20141022023202/http://www.nytimes.com/2011/02/17/science/17jeopardy-watson.html>) from the original on 22 October 2014. Retrieved 25 October 2014.
- Anadiotis, George (1 October 2020). "The state of AI in 2020: Democratization, industrialization, and the way to artificial general intelligence" (<https://www.zdnet.com/article/the-state-of-ai-in-2020-democratization-industrialization-and-the-way-to-artificial-general-intelligence/>). *ZDNet*. Retrieved 1 March 2021.

- Goertzel, Ben; Lian, Ruiting; Arel, Chen, Shuo (December 2010). "A world survey of artificial brain proj inspired cognitive architectures". *Neurocomputing*. 74 (1–3): 30–49. doi:10.1016/j.neucom.2010.08.012 (<https://doi.org/10.1016/j.neucom.2010.08.012>).
- Robinson, A. J.; Fallside, F. (1987), "The utility driven dynamic error propagation network.", *Technical Report CUED/F-INFENG/TR.1*, Cambridge University Engineering Department
- Hochreiter, Sepp (1991). *Untersuchungen zu dynamischen neuronalen Netzen* (<https://web.archive.org/web/20150306075401/http://people.idsia.ch/~juergen/SeppHochreiter1991ThesisAdvisorSchmidhuber.pdf>) (PDF) (diploma thesis). Munich: Institut f. Informatik, Technische Univ. Archived from the original (<http://people.idsia.ch/~juergen/SeppHochreiter1991ThesisAdvisorSchmidhuber.pdf>) (PDF) on 6 March 2015. Retrieved 16 April 2016.
- Williams, R. J.; Zipser, D. (1994), "Gradient-based learning algorithms for recurrent networks and their computational complexity", *Back-propagation: Theory, Architectures and Applications*, Hillsdale, NJ: Erlbaum
- Hochreiter, Sepp; Schmidhuber, Jürgen (1997), "Long Short-Term Memory", *Neural Computation*, 9 (8): 1735–1780, doi:10.1162/neco.1997.9.8.1735 (<https://doi.org/10.1162%2Fneco.1997.9.8.1735>), PMID 9377276 (<https://pubmed.ncbi.nlm.nih.gov/9377276/>), S2CID 1915014 (<https://api.semanticscholar.org/CorpusID:1915014>)
- Goodfellow, Ian; Bengio, Yoshua; Courville, Aaron (2016), *Deep Learning* (<https://web.archive.org/web/20160416111010/http://www.deeplearningbook.org/>), MIT Press., archived from the original (<http://www.deeplearningbook.org/>) on 16 April 2016, retrieved 12 November 2017
- Hinton, G.; Deng, L.; Yu, D.; Dahl, G.; Mohamed, A.; Jaitly, N.; Senior, A.; Vanhoucke, V.; Nguyen, P.; Sainath, T.; Kingsbury, B. (2012). "Deep Neural Networks for Acoustic Modeling in Speech Recognition – The shared views of four research groups". *IEEE Signal Processing Magazine*. 29 (6): 82–97. Bibcode:2012ISPM...29...82H (<https://ui.adsabs.harvard.edu/abs/2012ISPM...29...82H>). doi:10.1109/msp.2012.2205597 (<https://doi.org/10.1109%2Fmsp.2012.2205597>). S2CID 206485943 (<https://api.semanticscholar.org/CorpusID:206485943>).
- Schmidhuber, J. (2015). "Deep Learning in Neural Networks: An Overview". *Neural Networks*. 61: 85–117. arXiv:1404.7828 (<https://arxiv.org/abs/1404.7828>). doi:10.1016/j.neunet.2014.09.003 (<https://doi.org/10.1016%2Fj.neunet.2014.09.003>). PMID 25462637 (<https://pubmed.ncbi.nlm.nih.gov/25462637/>). S2CID 11715509 (<https://api.semanticscholar.org/CorpusID:11715509>).
- Linnainmaa, Seppo (1970). *The representation of the cumulative rounding error of an algorithm as a Taylor expansion of the local rounding errors* (Thesis) (in Finnish). Univ. Helsinki, 6–7.]
- Griewank, Andreas (2012). "Who Invented the Reverse Mode of Differentiation? Optimization Stories". *Documenta Mathematica, Extra Volume ISMP*: 389–400.
- Werbos, Paul (1974). *Beyond Regression: New Tools for Prediction and Analysis in the Behavioral Sciences* (Ph.D. thesis). Harvard University.
- Werbos, Paul (1982). *"Beyond Regression: New Tools for Prediction and Analysis in the Behavioral Sciences"* (<https://web.archive.org/web/20160414055503/http://werbos.com/NeuralSensitivityIFIPSeptember1981.pdf>) (PDF). *System Modeling and Optimization. Applications of advances in nonlinear sensitivity analysis*. Berlin, Heidelberg: Springer. Archived from the original (<http://werbos.com/NeuralSensitivityIFIPSeptember1981.pdf>) (PDF) on 14 April 2016. Retrieved 16 April 2016.
- "What is 'fuzzy logic'? Are there computers that are inherently fuzzy and do not apply the usual binary logic?" (<https://www.scientificamerican.com/article/what-is-fuzzy-logic-are-t/>). *Scientific American*. 21 October 1999. Retrieved 5 May 2018.
- Merkle, Daniel; Middendorf, Martin (2013). "Swarm Intelligence". In Burke, Edmund K.; Kendall, Graham (eds.). *Search Methodologies: Introductory Tutorials in Optimization and Decision Support Techniques*. Springer Science & Business Media. ISBN 978-1-4614-6940-7.

- van der Walt, Christiaan; Bernard "The characteristics that determine classifier performance" (https://web.archive.org/web/20090325194051/http://www.patternrecognition.co.za/publications/cvdwalt_data_characteristics_classifiers.pdf) (PDF). Archived from the original (http://www.patternrecognition.co.za/publications/cvdwalt_data_characteristics_classifiers.pdf) (PDF) on 25 March 2009. Retrieved 5 August 2009.
- Hutter, Marcus (2005). *Universal Artificial Intelligence*. Berlin: Springer. ISBN 978-3-540-22139-5.
- Howe, J. (November 1994). "Artificial Intelligence at Edinburgh University: a Perspective" (<http://www.inf.ed.ac.uk/about/AIhistory.html>). Archived (<https://web.archive.org/web/20070515072641/http://www.inf.ed.ac.uk/about/AIhistory.html>) from the original on 15 May 2007. Retrieved 30 August 2007.
- Galvan, Jill (1 January 1997). "Entering the Posthuman Collective in Philip K. Dick's "Do Androids Dream of Electric Sheep?". *Science Fiction Studies*. 24 (3): 413–429. JSTOR 4240644 (<https://www.jstor.org/stable/4240644>).
- McCauley, Lee (2007). "AI armageddon and the three laws of robotics". *Ethics and Information Technology*. 9 (2): 153–164. CiteSeerX 10.1.1.85.8904 (<https://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.85.8904>). doi:10.1007/s10676-007-9138-2 (<https://doi.org/10.1007%2Fs10676-007-9138-2>). S2CID 37272949 (<https://api.semanticscholar.org/CorpusID:37272949>).
- Buttazzo, G. (July 2001). "Artificial consciousness: Utopia or real possibility?". *Computer*. 34 (7): 24–30. doi:10.1109/2.933500 (<https://doi.org/10.1109%2F2.933500>).
- Anderson, Susan Leigh (2008). "Asimov's "three laws of robotics" and machine metaethics". *AI & Society*. 22 (4): 477–493. doi:10.1007/s00146-007-0094-5 (<https://doi.org/10.1007%2Fs00146-007-0094-5>). S2CID 1809459 (<https://api.semanticscholar.org/CorpusID:1809459>).
- Yudkowsky, E (2008), "Artificial Intelligence as a Positive and Negative Factor in Global Risk" (<http://intelligence.org/files/AIPosNegFactor.pdf>) (PDF), *Global Catastrophic Risks*, Oxford University Press, 2008, Bibcode:2008gcr..book..303Y (<https://ui.adsabs.harvard.edu/abs/2008gcr..book..303Y>)
- McGaughey, E (2018), *Will Robots Automate Your Job Away? Full Employment, Basic Income, and Economic Democracy* (https://web.archive.org/web/20180524201340/https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3044448), p. SSRN part 2(3), SSRN 3044448 (https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3044448), archived from the original (https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3044448) on 24 May 2018, retrieved 12 January 2018
- IGM Chicago (30 June 2017). "Robots and Artificial Intelligence" (<http://www.igmchicago.org/surveys/robots-and-artificial-intelligence>). www.igmchicago.org. Archived (<https://web.archive.org/web/20190501114826/http://www.igmchicago.org/surveys/robots-and-artificial-intelligence>) from the original on 1 May 2019. Retrieved 3 July 2019.
- Lohr, Steve (2017). "Robots Will Take Jobs, but Not as Fast as Some Fear, New Report Says" (<https://www.nytimes.com/2017/01/12/technology/robots-will-take-jobs-but-not-as-fast-as-some-fear-new-report-says.html>). *The New York Times*. Archived (<https://web.archive.org/web/20180114073704/https://www.nytimes.com/2017/01/12/technology/robots-will-take-jobs-but-not-as-fast-as-some-fear-new-report-says.html>) from the original on 14 January 2018. Retrieved 13 January 2018.
- Frey, Carl Benedikt; Osborne, Michael A (1 January 2017). "The future of employment: How susceptible are jobs to computerisation?". *Technological Forecasting and Social Change*. 114: 254–280. CiteSeerX 10.1.1.395.416 (<https://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.395.416>). doi:10.1016/j.techfore.2016.08.019 (<https://doi.org/10.1016%2Fj.techfore.2016.08.019>). ISSN 0040-1625 (<https://www.worldcat.org/issn/0040-1625>).
- Arntz, Melanie; Gregory, Terry; Zierahn, Ulrich (2016), "The risk of automation for jobs in OECD countries: A comparative analysis", *OECD Social, Employment, and Migration Working Papers* 189
- Morgenstern, Michael (9 May 2015). "Automation and anxiety" (<https://www.economist.com/news/special-report/21700758-will-smarter-machines-cause-mass-unemployment-automation-an-d-anxiety>). *The Economist*. Archived (<https://web.archive.org/web/20180112214621/https://www.economist.com/news/special-report/21700758-will-smarter-machines-cause-mass-unemployment-automation-an-d-anxiety>)

w.economist.com/news/special-report-automation-and-anxiety) from the original on 2018.

Artificial-machines-cause-mass-unemployment January 2018. Retrieved 13 January 2018.

- Mahdawi, Arwa (26 June 2017). "What jobs will still be around in 20 years? Read this to prepare your future" (<https://www.theguardian.com/us-news/2017/jun/26/jobs-future-automation-robots-skills-creative-health>). *The Guardian*. Archived (<https://web.archive.org/web/20180114021804/https://www.theguardian.com/us-news/2017/jun/26/jobs-future-automation-robots-skills-creative-health>) from the original on 14 January 2018. Retrieved 13 January 2018.
- Rubin, Charles (Spring 2003). "Artificial Intelligence and Human Nature" (<https://web.archive.org/web/20120611115223/http://www.thenewatlantis.com/publications/artificial-intelligence-and-human-nature>). *The New Atlantis*. 1: 88–100. Archived from the original (<http://www.thenewatlantis.com/publications/artificial-intelligence-and-human-nature>) on 11 June 2012.
- Bostrom, Nick (2014). *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press.
- Brooks, Rodney (10 November 2014). "artificial intelligence is a tool, not a threat" (<https://web.archive.org/web/20141112130954/http://www.rethinkrobotics.com/artificial-intelligence-tool-threat/>). Archived from the original (<http://www.rethinkrobotics.com/artificial-intelligence-tool-threat/>) on 12 November 2014.
- Sainato, Michael (19 August 2015). "Stephen Hawking, Elon Musk, and Bill Gates Warn About Artificial Intelligence" (<https://observer.com/2015/08/stephen-hawking-elon-musk-and-bill-gates-warn-about-artificial-intelligence/>). *Observer*. Archived (<https://web.archive.org/web/20151030053323/http://observer.com/2015/08/stephen-hawking-elon-musk-and-bill-gates-warn-about-artificial-intelligence/>) from the original on 30 October 2015. Retrieved 30 October 2015.
- Harari, Yuval Noah (October 2018). "Why Technology Favors Tyranny" (<https://www.theatlantic.com/magazine/archive/2018/10/yuval-noah-harari-technology-tyranny/568330>). *The Atlantic*.
- Robitzski, Dan (5 September 2018). "Five experts share what scares them the most about AI" (<https://futurism.com/artificial-intelligence-experts-fear/amp>). Archived (<https://web.archive.org/web/20191208094101/https://futurism.com/artificial-intelligence-experts-fear/amp>) from the original on 8 December 2019. Retrieved 8 December 2019.
- Goffrey, Andrew (2008). "Algorithm". In Fuller, Matthew (ed.). *Software studies: a lexicon* (https://archive.org/details/softwarestudiesl00full_007). Cambridge, Mass.: MIT Press. pp. 15 (https://archive.org/details/softwarestudiesl00full_007/page/n29)–20. ISBN 978-1-4356-4787-9.
- Lipartito, Kenneth (6 January 2011). *The Narrative and the Algorithm: Genres of Credit Reporting from the Nineteenth Century to Today* (https://mpra.ub.uni-muenchen.de/28142/1/MPRA_paper_28142.pdf) (PDF) (Unpublished manuscript), doi:10.2139/ssrn.1736283 (<https://doi.org/10.2139/ssrn.1736283>), S2CID 166742927 (<https://api.semanticscholar.org/CorpusID:166742927>)
- Goodman, Bryce; Flaxman, Seth (2017). "EU regulations on algorithmic decision-making and a "right to explanation" ". *AI Magazine*. 38 (3): 50. arXiv:1606.08813 (<https://arxiv.org/abs/1606.08813>). doi:10.1609/aimag.v38i3.2741 (<https://doi.org/10.1609/aimag.v38i3.2741>). S2CID 7373959 (<https://api.semanticscholar.org/CorpusID:7373959>).
- CNA (12 January 2019). "Commentary: Bad news. Artificial intelligence is biased" (<https://www.channelnewsasia.com/news/commentary/artificial-intelligence-big-data-bias-hiring-loans-key-challenge-11097374>). CNA. Archived (<https://web.archive.org/web/20190112104421/https://www.channelnewsasia.com/news/commentary/artificial-intelligence-big-data-bias-hiring-loans-key-challenge-11097374>) from the original on 12 January 2019. Retrieved 19 June 2020.
- Larson, Jeff; Angwin, Julia (23 May 2016). "How We Analyzed the COMPAS Recidivism Algorithm" (<https://www.propublica.org/article/how-we-analyzed-the-compas-recidivism-algorithm>). ProPublica. Archived (<https://web.archive.org/web/20190429190950/https://www.propublica.org/article/how-we-analyzed-the-compas-recidivism-algorithm>) from the original on 29 April 2019. Retrieved 19 June 2020.
- Müller, Vincent C.; Bostrom, Nick (2014). "Future Progress in Artificial Intelligence: A Poll Among Experts" (http://www.sophia.de/pdf/2014_PT-AI_polls.pdf) (PDF). *AI Matters*. 1 (1): 9–11. doi:10.1145/2639475.2639478 (<https://doi.org/10.1145/2639475.2639478>).

S2CID 8510016 (<https://api.semanticscholar.org/web/20160115114604/ht>)



'D:8510016). Archived (https://web.archive.org/web/2014_01/01/D:8510016). Archived (https://web.archive.org/web/2014_01/01/D:8510016) (PDF) from the original on 15 January 2016.

- Cellan-Jones, Rory (2 December 2014). "Stephen Hawking warns artificial intelligence could end mankind" (<https://www.bbc.com/news/technology-30290540>). BBC News. Archived (<https://web.archive.org/web/20151030054329/http://www.bbc.com/news/technology-30290540>) from the original on 30 October 2015. Retrieved 30 October 2015.
- Rawlinson, Kevin (29 January 2015). "Microsoft's Bill Gates insists AI is a threat" (<https://www.bbc.co.uk/news/31047780>). BBC News. Archived (<https://web.archive.org/web/20150129183607/http://www.bbc.co.uk/news/31047780>) from the original on 29 January 2015. Retrieved 30 January 2015.
- Holley, Peter (28 January 2015). "Bill Gates on dangers of artificial intelligence: 'I don't understand why some people are not concerned'" (<https://www.washingtonpost.com/news/the-switch/wp/2015/01/28/bill-gates-on-dangers-of-artificial-intelligence-dont-understand-why-some-people-are-not-concerned/>). The Washington Post. ISSN 0190-8286 (<https://www.worldcat.org/issn/0190-8286>). Archived (<https://web.archive.org/web/20151030054330/https://www.washingtonpost.com/news/the-switch/wp/2015/01/28/bill-gates-on-dangers-of-artificial-intelligence-dont-understand-why-some-people-are-not-concerned/>) from the original on 30 October 2015. Retrieved 30 October 2015.
- Gibbs, Samuel (27 October 2014). "Elon Musk: artificial intelligence is our biggest existential threat" (<https://www.theguardian.com/technology/2014/oct/27/elon-musk-artificial-intelligence-a-i-biggest-existential-threat>). The Guardian. Archived (<https://web.archive.org/web/20151030054330/http://www.theguardian.com/technology/2014/oct/27/elon-musk-artificial-intelligence-a-i-biggest-existential-threat>) from the original on 30 October 2015. Retrieved 30 October 2015.
- Churm, Philip Andrew (14 May 2019). "Yuval Noah Harari talks politics, technology and migration" (<https://www.euronews.com/2019/05/14/a-i-is-as-threatening-as-climate-change-and-nuclear-war-says-philosopher-yuval-noah-harari>). euronews. Archived (<https://web.archive.org/web/20190514192815/https://www.euronews.com/2019/05/14/a-i-is-as-threatening-as-climate-change-and-nuclear-war-says-philosopher-yuval-noah-harari>) from the original on 14 May 2019. Retrieved 15 November 2020.
- Bostrom, Nick (2015). "What happens when our computers get smarter than we are?" (https://www.ted.com/talks/nick_bostrom_what_happens_when_our_computers_get_smarter_than_we_are/transcript). TED (conference). Archived (https://web.archive.org/web/20200725005719/https://www.ted.com/talks/nick_bostrom_what_happens_when_our_computers_get_smarter_than_we_are/transcript) from the original on 25 July 2020. Retrieved 30 January 2020.
- Post, Washington (2015). "Tech titans like Elon Musk are spending \$1 billion to save you from terminators" (<https://www.chicagotribune.com/bluesky/technology/ct-tech-titans-against-terminators-20151214-story.html>). Chicago Tribune. Archived (<https://web.archive.org/web/20160607121118/http://www.chicagotribune.com/bluesky/technology/ct-tech-titans-against-terminators-20151214-story.html>) from the original on 7 June 2016.
- Del Prado, Guia Marie (9 October 2015). "The mysterious artificial intelligence company Elon Musk invested in is developing game-changing smart computers" (<http://www.techinsider.io/mysterious-artificial-intelligence-company-elon-musk-investment-2015-10>). Tech Insider. Archived (<https://web.archive.org/web/20151030165333/http://www.techinsider.io/mysterious-artificial-intelligence-company-elon-musk-investment-2015-10>) from the original on 30 October 2015. Retrieved 30 October 2015.
- FastCompany (15 January 2015). "Elon Musk Is Donating \$10M Of His Own Money To Artificial Intelligence Research" (<http://www.fastcompany.com/3041007/fast-feed/elon-musk-is-donating-10m-of-his-own-money-to-artificial-intelligence-research>). Fast Company. Archived (<https://web.archive.org/web/20151030202356/http://www.fastcompany.com/3041007/fast-feed/elon-musk-is-donating-10m-of-his-own-money-to-artificial-intelligence-research>) from the original on 30 October 2015. Retrieved 30 October 2015.

- Thibodeau, Patrick (25 March 2017) (<https://searcherp.techtarget.com/fe>). *SearchERP*. Archivea (<https://web.archive.org/web/20190506173749/https://searcherp.techtarget.com/news/252460208/Oracle-CEO-Mark-Hurd-sees-no-reason-to-fear-ERP-AI>) from the original on 6 May 2019. Retrieved 6 May 2019.
- Bhardwaj, Prachi (24 May 2018). "Mark Zuckerberg responds to Elon Musk's paranoia about AI: 'AI is going to... help keep our communities safe.' " (<https://www.businessinsider.com/mark-zuckerberg-shares-thoughts-elon-musks-ai-2018-5>). *Business Insider*. Archived (<https://web.archive.org/web/20190506173756/https://www.businessinsider.com/mark-zuckerberg-shares-thoughts-elon-musks-ai-2018-5>) from the original on 6 May 2019. Retrieved 6 May 2019.
- Geist, Edward Moore (9 August 2015). "Is artificial intelligence really an existential threat to humanity?" (<http://thebulletin.org/artificial-intelligence-really-existential-threat-humanity8577>). *Bulletin of the Atomic Scientists*. Archived (<https://web.archive.org/web/20151030054330/http://thebulletin.org/artificial-intelligence-really-existential-threat-humanity8577>) from the original on 30 October 2015. Retrieved 30 October 2015.
- Madrigal, Alexis C. (27 February 2015). "The case against killer robots, from a guy actually working on artificial intelligence" (<https://www.hrw.org/report/2012/11/19/losing-humanity/case-against-killer-robots>). *Fusion.net*. Archived (<https://web.archive.org/web/20160204175716/http://fusion.net/story/54583/the-case-against-killer-robots-from-a-guy-actually-building-ai/>) from the original on 4 February 2016. Retrieved 31 January 2016.
- Lee, Timothy B. (22 August 2014). "Will artificial intelligence destroy humanity? Here are 5 reasons not to worry" (<https://www.vox.com/2014/8/22/6043635/5-reasons-we-shouldnt-worry-about-super-intelligent-computers-taking>). *Vox*. Archived (<https://web.archive.org/web/20151030092203/http://www.vox.com/2014/8/22/6043635/5-reasons-we-shouldnt-worry-about-super-intelligent-computers-taking>) from the original on 30 October 2015. Retrieved 30 October 2015.
- Law Library of Congress (U.S.). Global Legal Research Directorate, issuing body. (2019). *Regulation of artificial intelligence in selected jurisdictions*. LCCN 2019668143 (<https://lccn.loc.gov/2019668143>). OCLC 1110727808 (<https://www.worldcat.org/oclc/1110727808>).
- UNESCO Science Report: the Race Against Time for Smarter Development (<https://unesdoc.unesco.org/ark:/48223/pf0000377433/PDF/377433eng.pdf.multi>). Paris: UNESCO. 11 June 2021. ISBN 978-92-3-100450-6.
- Berryhill, Jamie; Heang, Kévin Kok; Clogher, Rob; McBride, Keegan (2019). *Hello, World: Artificial Intelligence and its Use in the Public Sector* (<https://oecd-opsi.org/wp-content/uploads/2019/11/AI-Report-Online.pdf>) (PDF). Paris: OECD Observatory of Public Sector Innovation. Archived (<https://web.archive.org/web/20191220021331/https://oecd-opsi.org/wp-content/uploads/2019/11/AI-Report-Online.pdf>) (PDF) from the original on 20 December 2019. Retrieved 9 August 2020.
- Barfield, Woodrow; Pagallo, Ugo (2018). *Research handbook on the law of artificial intelligence*. Cheltenham, UK. ISBN 978-1-78643-904-8. OCLC 1039480085 (<https://www.worldcat.org/oclc/1039480085>).
- Iphofen, Ron; Kritikos, Mihalis (3 January 2019). "Regulating artificial intelligence and robotics: ethics by design in a digital society". *Contemporary Social Science*. 16 (2): 170–184. doi:10.1080/21582041.2018.1563803 (<https://doi.org/10.1080%2F21582041.2018.1563803>). ISSN 2158-2041 (<https://www.worldcat.org/issn/2158-2041>). S2CID 59298502 (<https://api.semanticscholar.org/CorpusID:59298502>).
- Wirtz, Bernd W.; Weyerer, Jan C.; Geyer, Carolin (24 July 2018). "Artificial Intelligence and the Public Sector – Applications and Challenges" (<https://zenodo.org/record/3569435>). *International Journal of Public Administration*. 42 (7): 596–615. doi:10.1080/01900692.2018.1498103 (<https://doi.org/10.1080%2F01900692.2018.1498103>). ISSN 0190-0692 (<https://www.worldcat.org/issn/0190-0692>). S2CID 158829602 (<https://api.semanticscholar.org/CorpusID:158829602>). Archived (<https://web.archive.org/web/20200818131415/https://zenodo.org/record/3569435>) from the original on 18 August 2020. Retrieved 22 August 2020.

- Buitenhuis, Miriam C (2019). "Towards a More Human Future: The Impact of Artificial Intelligence" (<https://doi.org/10.1017%2Ferr.2019.8>). *European Journal of Artificial Intelligence*. 10 (1): 41–59. doi:10.1017/err.2019.8 (<https://doi.org/10.1017%2Ferr.2019.8>). ISSN 1867-299X (<https://www.worldcat.org/issn/1867-299X>).
- Wallach, Wendell (2010). *Moral Machines*. Oxford University Press.
- Brown, Eileen (5 November 2019). "Half of Americans do not believe deepfake news could target them online" (<https://www.zdnet.com/article/half-of-americans-do-not-believe-deepfake-news-could-target-them-online/>). ZDNet. Archived (<https://web.archive.org/web/20191106035012/https://www.zdnet.com/article/half-of-americans-do-not-believe-deepfake-news-could-target-them-online/>) from the original on 6 November 2019. Retrieved 3 December 2019.
- Frangoul, Anmar (14 June 2019). "A Californian business is using A.I. to change the way we think about energy storage" (<https://www.cnbc.com/2019/06/14/the-business-using-ai-to-change-how-we-think-about-energy-storage.html>). CNBC. Archived (<https://web.archive.org/web/20200725044735/https://www.cnbc.com/2019/06/14/the-business-using-ai-to-change-how-we-think-about-energy-storage.html>) from the original on 25 July 2020. Retrieved 5 November 2019.
- "The Economist Explains: Why firms are piling into artificial intelligence" (<https://www.economist.com/blogs/economist-explains/2016/04/economist-explains>). *The Economist*. 31 March 2016. Archived (<https://web.archive.org/web/20160508010311/http://www.economist.com/blogs/economist-explains/2016/04/economist-explains>) from the original on 8 May 2016. Retrieved 19 May 2016.
- Lohr, Steve (28 February 2016). "The Promise of Artificial Intelligence Unfolds in Small Steps" (<https://www.nytimes.com/2016/02/29/technology/the-promise-of-artificial-intelligence-unfolds-in-small-steps.html>). *The New York Times*. Archived (<https://web.archive.org/web/20160229171843/http://www.nytimes.com/2016/02/29/technology/the-promise-of-artificial-intelligence-unfolds-in-small-steps.html>) from the original on 29 February 2016. Retrieved 29 February 2016.
- Smith, Mark (22 July 2016). "So you think you chose to read this article?" (<https://www.bbc.co.uk/news/business-36837824>). BBC News. Archived (<https://web.archive.org/web/20160725205007/http://www.bbc.co.uk/news/business-36837824>) from the original on 25 July 2016.
- Aletras, N.; Tsarapatsanis, D.; Preotiuc-Pietro, D.; Lampos, V. (2016). "Predicting judicial decisions of the European Court of Human Rights: a Natural Language Processing perspective" (<https://doi.org/10.7717%2Fpeerj-cs.93>). *PeerJ Computer Science*. 2: e93. doi:10.7717/peerj-cs.93 (<https://doi.org/10.7717%2Fpeerj-cs.93>).
- Cadena, Cesar; Carbone, Luca; Carrillo, Henry; Latif, Yasir; Scaramuzza, Davide; Neira, Jose; Reid, Ian; Leonard, John J. (December 2016). "Past, Present, and Future of Simultaneous Localization and Mapping: Toward the Robust-Perception Age". *IEEE Transactions on Robotics*. 32 (6): 1309–1332. arXiv:1606.05830 (<https://arxiv.org/abs/1606.05830>). doi:10.1109/TRO.2016.2624754 (<https://doi.org/10.1109%2FTRO.2016.2624754>). S2CID 2596787 (<https://api.semanticscholar.org/CorpusID:2596787>).
- Cambria, Erik; White, Bebo (May 2014). "Jumping NLP Curves: A Review of Natural Language Processing Research [Review Article]". *IEEE Computational Intelligence Magazine*. 9 (2): 48–57. doi:10.1109/MCI.2014.2307227 (<https://doi.org/10.1109%2FMCI.2014.2307227>). S2CID 206451986 (<https://api.semanticscholar.org/CorpusID:206451986>).
- Vincent, James (7 November 2019). "OpenAI has published the text-generating AI it said was too dangerous to share" (<https://www.theverge.com/2019/11/7/20953040/openai-text-generation-ai-gpt-2-full-model-release-1-5b-parameters>). The Verge. Archived (<https://web.archive.org/web/20200611054114/https://www.theverge.com/2019/11/7/20953040/openai-text-generation-ai-gpt-2-full-model-release-1-5b-parameters>) from the original on 11 June 2020. Retrieved 11 June 2020.
- Jordan, M. I.; Mitchell, T. M. (16 July 2015). "Machine learning: Trends, perspectives, and prospects". *Science*. 349 (6245): 255–260. Bibcode:2015Sci...349..255J (<https://ui.adsabs.harvard.edu/abs/2015Sci...349..255J>). doi:10.1126/science.aaa8415 (<https://doi.org/10.1126%2Fscience.aaa8415>). PMID 26185243 (<https://pubmed.ncbi.nlm.nih.gov/26185243>). S2CID 677218 (<https://api.semanticscholar.org/CorpusID:677218>).

- Maschafilm (2010). "Content: Plug and Pray Film". Archived from the original on 12 February 2016. [Intelligence – Robots –](http://www.pluganpray-tiim.ae/en/content.html) (http://www.pluganpray-tiim.ae/en/content.html) from the original on 12 February 2016.
- Evans, Woody (2015). "Posthuman Rights: Dimensions of Transhuman Worlds" (https://doi.org/10.5209/rev_TK.2015.v12.n2.49072). *Teknokultura*. 12 (2). doi:10.5209/rev_TK.2015.v12.n2.49072 (https://doi.org/10.5209/rev_TK.2015.v12.n2.49072).
- Waddell, Kaveh (2018). "Chatbots Have Entered the Uncanny Valley" (https://www.theatlantic.com/technology/archive/2017/04/uncanny-valley-digital-assistants/523806/). *The Atlantic*. Archived (https://web.archive.org/web/20180424202350/https://www.theatlantic.com/technology/archive/2017/04/uncanny-valley-digital-assistants/523806/) from the original on 24 April 2018. Retrieved 24 April 2018.
- Poria, Soujanya; Cambria, Erik; Bajpai, Rajiv; Hussain, Amir (September 2017). "A review of affective computing: From unimodal analysis to multimodal fusion" (http://researchrepository.napier.ac.uk/Output/1792429). *Information Fusion*. 37: 98–125. doi:10.1016/j.inffus.2017.02.003 (https://doi.org/10.1016%2Fj.inffus.2017.02.003). hdl:1893/25490 (https://hdl.handle.net/1893%2F25490).
- "Robots could demand legal rights" (http://news.bbc.co.uk/2/hi/technology/6200005.stm). *BBC News*. 21 December 2006. Archived (https://web.archive.org/web/20191015042628/http://news.bbc.co.uk/2/hi/technology/6200005.stm) from the original on 15 October 2019. Retrieved 3 February 2011.
- Horst, Steven (2005). "The Computational Theory of Mind" (http://plato.stanford.edu/entries/computational-mind). *The Stanford Encyclopedia of Philosophy*.
- Omohundro, Steve (2008). *The Nature of Self-Improving Artificial Intelligence*. presented and distributed at the 2007 Singularity Summit, San Francisco, CA.
- Ford, Martin; Colvin, Geoff (6 September 2015). "Will robots create more jobs than they destroy?" (https://www.theguardian.com/technology/2015/sep/06/will-robots-create-destroy-jobs). *The Guardian*. Archived (https://web.archive.org/web/20180616204119/https://www.theguardian.com/technology/2015/sep/06/will-robots-create-destroy-jobs) from the original on 16 June 2018. Retrieved 13 January 2018.
- *White Paper: On Artificial Intelligence – A European approach to excellence and trust* (https://ec.europa.eu/info/sites/info/files/commission-white-paper-artificial-intelligence-feb2020_en.pdf) (PDF). Brussels: European Commission. 2020. Archived (https://web.archive.org/web/20200220173419/https://ec.europa.eu/info/sites/info/files/commission-white-paper-artificial-intelligence-feb2020_en.pdf) (PDF) from the original on 20 February 2020. Retrieved 20 February 2020.
- Anderson, Michael; Anderson, Susan Leigh (2011). *Machine Ethics*. Cambridge University Press.
- "Machine Ethics" (https://web.archive.org/web/20141129044821/http://www.aaai.org/Library/Symposia/Fall/fs05-06). aaai.org. Archived from the original (http://www.aaai.org/Library/Symposia/Fall/fs05-06) on 29 November 2014.
- Russell, Stuart (8 October 2019). *Human Compatible: Artificial Intelligence and the Problem of Control*. United States: Viking. ISBN 978-0-525-55861-3. OCLC 1083694322 (https://www.worldcat.org/oclc/1083694322).
- "AI set to exceed human brain power" (http://www.cnn.com/2006/TECH/science/07/24/ai.bostrom/). CNN. 9 August 2006. Archived (https://web.archive.org/web/20080219001624/http://www.cnn.com/2006/TECH/science/07/24/ai.bostrom/) from the original on 19 February 2008.
- "Robots could demand legal rights" (http://news.bbc.co.uk/2/hi/technology/6200005.stm). *BBC News*. 21 December 2006. Archived (https://web.archive.org/web/20191015042628/http://news.bbc.co.uk/2/hi/technology/6200005.stm) from the original on 15 October 2019. Retrieved 3 February 2011.

- "Kismet" (<http://www.ai.mit.edu/projects/kismet/kismet.html>). MIT Artificial Intelligence Laboratory,    Archived (<https://web.archive.org/web/20141017040432/http://www.ai.mit.edu/projects/numanoid-robotics-group/kismet/kismet.html>) from the original on 17 October 2014. Retrieved 25 October 2014.
- Smolar, Stephen W.; Zhang, HongJiang (1994). "Content based video indexing and retrieval". *IEEE MultiMedia*. 1 (2): 62–72. doi:10.1109/93.311653 (<https://doi.org/10.1109%2F93.311653>). S2CID 32710913 (<https://api.semanticscholar.org/CorpusID:32710913>).
- Neumann, Bernd; Möller, Ralf (January 2008). "On scene interpretation with description logics". *Image and Vision Computing*. 26 (1): 82–101. doi:10.1016/j.imavis.2007.08.013 (<https://doi.org/10.1016%2Fj.imavis.2007.08.013>).
- Kuperman, G. J.; Reichley, R. M.; Bailey, T. C. (1 July 2006). "Using Commercial Knowledge Bases for Clinical Decision Support: Opportunities, Hurdles, and Recommendations" (<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC1513681>). *Journal of the American Medical Informatics Association*. 13 (4): 369–371. doi:10.1197/jamia.M2055 (<https://doi.org/10.1197%2Fjamia.M2055>). PMC 1513681 (<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC1513681>). PMID 16622160 (<https://pubmed.ncbi.nlm.nih.gov/16622160>).
- McGarry, Ken (1 December 2005). "A survey of interestingness measures for knowledge discovery". *The Knowledge Engineering Review*. 20 (1): 39–61. doi:10.1017/S0269888905000408 (<https://doi.org/10.1017%2FS0269888905000408>). S2CID 14987656 (<https://api.semanticscholar.org/CorpusID:14987656>).
- Bertini, M; Del Bimbo, A; Torniai, C (2006). "Automatic annotation and semantic retrieval of video sequences using multimedia ontologies". *MM '06 Proceedings of the 14th ACM international conference on Multimedia*. 14th ACM international conference on Multimedia. Santa Barbara: ACM. pp. 679–682.
- Kahneman, Daniel (25 October 2011). *Thinking, Fast and Slow* (<https://books.google.com/books?id=ZuKTvERuPG8C>). Macmillan. ISBN 978-1-4299-6935-2. Retrieved 8 April 2012.
- Turing, Alan (1948), "Machine Intelligence", in Copeland, B. Jack (ed.), *The Essential Turing: The ideas that gave birth to the computer age*, Oxford: Oxford University Press, p. 412, ISBN 978-0-19-825080-7
- Domingos, Pedro (22 September 2015). *The Master Algorithm: How the Quest for the Ultimate Learning Machine Will Remake Our World*. Basic Books. ISBN 978-0465065707.
- Minsky, Marvin (1986), *The Society of Mind*, Simon and Schuster
- Pinker, Steven (4 September 2007) [1994], *The Language Instinct*, Perennial Modern Classics, Harper, ISBN 978-0-06-133646-1
- Chalmers, David (1995). "Facing up to the problem of consciousness" (<http://www.imprint.co.uk/chalmers.html>). *Journal of Consciousness Studies*. 2 (3): 200–219. Archived (<https://web.archive.org/web/20050308163649/http://www.imprint.co.uk/chalmers.html>) from the original on 8 March 2005. Retrieved 11 October 2018.
- Roberts, Jacob (2016). "Thinking Machines: The Search for Artificial Intelligence" (<https://web.archive.org/web/20180819152455/https://www.sciencehistory.org/distillations/magazine/thinking-machines-the-search-for-artificial-intelligence>). *Distillations*. Vol. 2, no. 2. pp. 14–23. Archived from the original (<https://www.sciencehistory.org/distillations/magazine/thinking-machines-the-search-for-artificial-intelligence>) on 19 August 2018. Retrieved 20 March 2018.
- Pennachin, C.; Goertzel, B. (2007). "Contemporary Approaches to Artificial General Intelligence". *Artificial General Intelligence*. Cognitive Technologies. Berlin, Heidelberg: Springer. doi:10.1007/978-3-540-68677-4_1 (https://doi.org/10.1007%2F978-3-540-68677-4_1). ISBN 978-3-540-23733-4.
- "Ask the AI experts: What's driving today's progress in AI?" (<https://www.mckinsey.com/business-functions/mckinsey-analytics/our-insights/ask-the-ai-experts-whats-driving-todays-progress-in-ai>). McKinsey & Company. Archived (<https://web.archive.org/web/20180413190018/https://www.mckinsey.com/business-functions/mckinsey-analytics/our-insights/ask-the-ai-experts-whats-driving-todays-progress-in-ai>) from the original on 13 April 2018. Retrieved 13 April 2018.

- Ransbotham, Sam; Kiron, David; S, Martin (6 September 2017). "Reshaping Business With Artificial Intelligence". *MIT Sloan Management Review*. Archived (<https://web.archive.org/web/20180519171905/https://sloanreview.mit.edu/projects/reshaping-business-with-artificial-intelligence/>) from the original on 19 May 2018. Retrieved 2 May 2018.
- Lorica, Ben (18 December 2017). "The state of AI adoption" (<https://www.oreilly.com/ideas/the-state-of-ai-adoption>). O'Reilly Media. Archived (<https://web.archive.org/web/20180502140700/https://www.oreilly.com/ideas/the-state-of-ai-adoption>) from the original on 2 May 2018. Retrieved 2 May 2018.
- "AlphaGo – Google DeepMind" (<https://wildoftech.com/alphago-google-deepmind>). Archived (<https://web.archive.org/web/20160310191926/https://wildoftech.com/alphago-google-deepmind/>) from the original on 20 October 2021.
- Asada, M.; Hosoda, K.; Kuniyoshi, Y.; Ishiguro, H.; Inui, T.; Yoshikawa, Y.; Ogino, M.; Yoshida, C. (2009). "Cognitive developmental robotics: a survey". *IEEE Transactions on Autonomous Mental Development*. 1 (1): 12–34. doi:10.1109/tamd.2009.2021702 (<https://doi.org/10.1109/tamd.2009.2021702>). S2CID 10168773 (<https://api.semanticscholar.org/CorpusID:10168773>).
- Ashok83 (10 September 2019). "How AI Is Getting Groundbreaking Changes In Talent Management And HR Tech" (<https://hackernoon.com/how-ai-is-getting-groundbreaking-changes-in-talent-management-and-hr-tech-d24ty3zzd>). Hackernoon. Archived (<https://web.archive.org/web/20190911083500/https://hackernoon.com/how-ai-is-getting-groundbreaking-changes-in-talent-management-and-hr-tech-d24ty3zzd>) from the original on 11 September 2019. Retrieved 14 February 2020.
- Berlinski, David (2000). *The Advent of the Algorithm* (<https://archive.org/details/adventofalgorith0000berl>). Harcourt Books. ISBN 978-0-15-601391-8. OCLC 46890682 (<https://www.worldcat.org/oclc/46890682>). Archived (<https://web.archive.org/web/20200726215744/https://archive.org/details/adventofalgorith0000berl>) from the original on 26 July 2020. Retrieved 22 August 2020.
- Brooks, Rodney (1990). "Elephants Don't Play Chess" (<http://people.csail.mit.edu/brooks/papers/elephants.pdf>) (PDF). *Robotics and Autonomous Systems*. 6 (1–2): 3–15. CiteSeerX 10.1.1.588.7539 (<https://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.588.7539>). doi:10.1016/S0921-8890(05)80025-9 ([https://doi.org/10.1016/S0921-8890\(05\)80025-9](https://doi.org/10.1016/S0921-8890(05)80025-9)). Archived (<https://web.archive.org/web/20070809020912/http://people.csail.mit.edu/brooks/papers/elephants.pdf>) (PDF) from the original on 9 August 2007.
- Butler, Samuel (13 June 1863). "Darwin among the Machines" (<http://www.nzetc.org/tm/scholarly/tei-ButFir-t1-g1-t1-g1-t4-body.html>). Letters to the Editor. *The Press*. Christchurch, New Zealand. Archived (<https://web.archive.org/web/20080919172551/http://www.nzetc.org/tm/scholarly/tei-ButFir-t1-g1-t1-g1-t4-body.html>) from the original on 19 September 2008. Retrieved 16 October 2014 – via Victoria University of Wellington.
- Clark, Jack (2015a). "Musk-Backed Group Probes Risks Behind Artificial Intelligence" (<https://www.bloomberg.com/news/articles/2015-07-01/musk-backed-group-probes-risks-behind-artificial-intelligence>). Bloomberg.com. Archived (<https://web.archive.org/web/20151030202356/http://www.bloomberg.com/news/articles/2015-07-01/musk-backed-group-probes-risks-behind-artificial-intelligence>) from the original on 30 October 2015. Retrieved 30 October 2015.
- Clark, Jack (2015b). "Why 2015 Was a Breakthrough Year in Artificial Intelligence" (<https://www.bloomberg.com/news/articles/2015-12-08/why-2015-was-a-breakthrough-year-in-artificial-intelligence>). Bloomberg.com. Archived (<https://web.archive.org/web/20161123053855/https://www.bloomberg.com/news/articles/2015-12-08/why-2015-was-a-breakthrough-year-in-artificial-intelligence>) from the original on 23 November 2016. Retrieved 23 November 2016.
- Dennett, Daniel (1991). *Consciousness Explained*. The Penguin Press. ISBN 978-0-7139-9037-9.
- Dreyfus, Hubert (1972). *What Computers Can't Do*. New York: MIT Press. ISBN 978-0-06-011082-6.

- Dreyfus, Hubert; Dreyfus, Stuart (1992). *Mind over Machine: The Power of Human Intuition and Expertise in the Era of the Computer*. Oxford, UK: Blackwell. ISBN 978-0-631-19080-6. Archived (<https://web.archive.org/web/200726131414/https://archive.org/details/mindovermachinep00drey>) from the original on 26 July 2020. Retrieved 22 August 2020.
- Dyson, George (1998). *Darwin among the Machines* (<https://archive.org/details/darwinamongmachi00dys0>). Allan Lane Science. ISBN 978-0-7382-0030-9. Archived (<https://web.archive.org/web/20200726131443/https://archive.org/details/darwinamongmachi00dys0>) from the original on 26 July 2020. Retrieved 22 August 2020.
- Edelson, Edward (1991). *The Nervous System* (<https://archive.org/details/nervoussystem0000edel>). New York: Chelsea House. ISBN 978-0-7910-0464-7. Archived (<https://web.archive.org/web/20200726131758/https://archive.org/details/nervoussystem0000edel>) from the original on 26 July 2020. Retrieved 18 November 2019.
- Fearn, Nicholas (2007). *The Latest Answers to the Oldest Questions: A Philosophical Adventure with the World's Greatest Thinkers*. New York: Grove Press. ISBN 978-0-8021-1839-4.
- Haugeland, John (1985). *Artificial Intelligence: The Very Idea*. Cambridge, Mass.: MIT Press. ISBN 978-0-262-08153-5.
- Hawkins, Jeff; Blakeslee, Sandra (2005). *On Intelligence*. New York: Owl Books. ISBN 978-0-8050-7853-4.
- Henderson, Mark (24 April 2007). "Human rights for robots? We're getting carried away" (<http://www.thetimes.co.uk/tto/technology/article1966391.ece>). *The Times Online*, London. Archived (<https://web.archive.org/web/20140531104850/http://www.thetimes.co.uk/tto/technology/article1966391.ece>) from the original on 31 May 2014. Retrieved 31 May 2014.
- Kahneman, Daniel; Slovic, D.; Tversky, Amos (1982). *Judgment under uncertainty: Heuristics and biases*. *Science*. Vol. 185. New York: Cambridge University Press. pp. 1124–1131. Bibcode:1974Sci...185.1124T (<https://ui.adsabs.harvard.edu/abs/1974Sci...185.1124T>). doi:10.1126/science.185.4157.1124 (<https://doi.org/10.1126%2Fscience.185.4157.1124>). ISBN 978-0-521-28414-1. PMID 17835457 (<https://pubmed.ncbi.nlm.nih.gov/17835457/>). S2CID 143452957 (<https://api.semanticscholar.org/CorpusID:143452957>).
- Katz, Yarden (1 November 2012). "Noam Chomsky on Where Artificial Intelligence Went Wrong" (https://www.theatlantic.com/technology/archive/2012/11/noam-chomsky-on-where-artificial-intelligence-went-wrong/261637/?single_page=true). *The Atlantic*. Archived (https://web.archive.org/web/20190228154403/https://www.theatlantic.com/technology/archive/2012/11/noam-chomsky-on-where-artificial-intelligence-went-wrong/261637/?single_page=true) from the original on 28 February 2019. Retrieved 26 October 2014.
- Kurzweil, Ray (2005). *The Singularity is Near*. Penguin Books. ISBN 978-0-670-03384-3.
- Langley, Pat (2011). "The changing science of machine learning" (<https://doi.org/10.1007%2Fs10994-011-5242-y>). *Machine Learning*. 82 (3): 275–279. doi:10.1007/s10994-011-5242-y (<https://doi.org/10.1007%2Fs10994-011-5242-y>).
- Legg, Shane; Hutter, Marcus (15 June 2007). "A Collection of Definitions of Intelligence". arXiv:0706.3639 (<https://arxiv.org/abs/0706.3639>) [cs.AI] (<https://arxiv.org/archive/cs.AI>).
- Lenat, Douglas; Guha, R. V. (1989). *Building Large Knowledge-Based Systems*. Addison-Wesley. ISBN 978-0-201-51752-1.
- Lighthill, James (1973). "Artificial Intelligence: A General Survey". *Artificial Intelligence: a paper symposium*. Science Research Council.
- Lombardo, P; Boehm, I; Nairz, K (2020). "RadioComics – Santa Claus and the future of radiology" (<https://doi.org/10.1016%2Fj.ejrad.2019.108771>). *Eur J Radiol*. 122 (1): 108771. doi:10.1016/j.ejrad.2019.108771 (<https://doi.org/10.1016%2Fj.ejrad.2019.108771>). PMID 31835078 (<https://pubmed.ncbi.nlm.nih.gov/31835078/>).
- Lungarella, M.; Metta, G.; Pfeifer, R.; Sandini, G. (2003). "Developmental robotics: a survey". *Connection Science*. 15 (4): 151–190. CiteSeerX 10.1.1.83.7615 (<https://citeseerx.ist.psu.edu/>

- [viewdoc/summary?doi=10.1.1.83.7](#) 40090310001655110 (<https://doi.org/10.1080%2F0954009031000165110>)
- 40090310001655110 (<https://api.semanticscholar.org/CorpusID:1452734>).
■ Maker, Meg Houston (2006). "AI@50: AI Past, Present, Future" (https://web.archive.org/web/20070103222615/http://www.engagingexperience.com/2006/07/ai50_ai_past_pr.html). Dartmouth College. Archived from the original (http://www.engagingexperience.com/2006/07/ai50_ai_past_pr.html) on 3 January 2007. Retrieved 16 October 2008.
- McCarthy, John; Minsky, Marvin; Rochester, Nathan; Shannon, Claude (1955). "A Proposal for the Dartmouth Summer Research Project on Artificial Intelligence" (<https://web.archive.org/web/20070826230310/http://www-formal.stanford.edu/jmc/history/dartmouth/dartmouth.html>). Archived from the original (<http://www-formal.stanford.edu/jmc/history/dartmouth/dartmouth.html>) on 26 August 2007. Retrieved 30 August 2007.
- Minsky, Marvin (1967). *Computation: Finite and Infinite Machines* (<https://archive.org/details/computationfinit0000mins>). Englewood Cliffs, N.J.: Prentice-Hall. ISBN 978-0-13-165449-5. Archived (<https://web.archive.org/web/20200726131743/https://archive.org/details/computationfinit0000mins>) from the original on 26 July 2020. Retrieved 18 November 2019.
- Moravec, Hans (1988). *Mind Children* (<https://archive.org/details/mindchildrenfutu00mora>). Harvard University Press. ISBN 978-0-674-57616-2. Archived (<https://web.archive.org/web/20200726131644/https://archive.org/details/mindchildrenfutu00mora>) from the original on 26 July 2020. Retrieved 18 November 2019.
- NRC (United States National Research Council) (1999). "Developments in Artificial Intelligence". *Funding a Revolution: Government Support for Computing Research*. National Academy Press.
- Newell, Allen; Simon, H. A. (1976). "Computer Science as Empirical Inquiry: Symbols and Search" (<https://doi.org/10.1145%2F360018.360022>). *Communications of the ACM*. 19 (3): 113–126. doi:10.1145/360018.360022 (<https://doi.org/10.1145%2F360018.360022>)..
- Nilsson, Nils (1983). "Artificial Intelligence Prepares for 2001" (<https://ai.stanford.edu/~nilsson/OnlinePubs-Nils/General%20Essays/AIMag04-04-002.pdf>) (PDF). *AI Magazine*. 1 (1). Archived (<https://web.archive.org/web/20200817194457/http://ai.stanford.edu/~nilsson/OnlinePubs-Nils/General%20Essays/AIMag04-04-002.pdf>) (PDF) from the original on 17 August 2020. Retrieved 22 August 2020. Presidential Address to the Association for the Advancement of Artificial Intelligence.
- Oudeyer, P-Y. (2010). "On the impact of robotics in behavioral and cognitive sciences: from insect navigation to human cognitive development" (<http://www.pyoudeyer.com/IEEETAMDOudeyer10.pdf>) (PDF). *IEEE Transactions on Autonomous Mental Development*. 2 (1): 2–16. doi:10.1109/tamd.2009.2039057 (<https://doi.org/10.1109%2Ftamd.2009.2039057>). S2CID 6362217 (<https://api.semanticscholar.org/CorpusID:6362217>). Archived (<https://web.archive.org/web/20181003202543/http://www.pyoudeyer.com/IEEETAMDOudeyer10.pdf>) (PDF) from the original on 3 October 2018. Retrieved 4 June 2013.
- Schank, Roger C. (1991). "Where's the AI". *AI magazine*. Vol. 12, no. 4.
- Searle, John (1980). "Minds, Brains and Programs" (<http://cogprints.org/7150/1/10.1.1.83.5248.pdf>) (PDF). *Behavioral and Brain Sciences*. 3 (3): 417–457. doi:10.1017/S0140525X00005756 (<https://doi.org/10.1017%2FS0140525X00005756>). S2CID 55303721 (<https://api.semanticscholar.org/CorpusID:55303721>). Archived (<https://web.archive.org/web/20190317230215/http://cogprints.org/7150/1/10.1.1.83.5248.pdf>) (PDF) from the original on 17 March 2019. Retrieved 22 August 2020.
- Searle, John (1999). *Mind, language and society* (<https://archive.org/details/mindlanguagesoci00sear>). New York: Basic Books. ISBN 978-0-465-04521-1. OCLC 231867665 (<https://www.worldcat.org/oclc/231867665>). Archived (<https://web.archive.org/web/20200726220615/https://archive.org/details/mindlanguagesoci00sear>) from the original on 26 July 2020. Retrieved 22 August 2020.
- Simon, H. A. (1965). *The Shape of Automation for Men and Management* (<https://archive.org/details/shapeofautomatio00simo>). New York: Harper & Row. Archived (<https://web.archive.org/w>

- [eb/20200726131655/https://archive.org/details/tomatito00simo](https://web.archive.org/web/20200726131655/https://archive.org/details/tomatito00simo)) from the original on 26 July 2020. Retrieved 18 November 2020.
- Solomonoff, Ray (1956). *An Inductive Inference Machine* (<http://world.std.com/~rjs/indinf56.pdf>) (PDF). Dartmouth Summer Research Conference on Artificial Intelligence. Archived (<https://web.archive.org/web/20110426161749/http://world.std.com/~rjs/indinf56.pdf>) (PDF) from the original on 26 April 2011. Retrieved 22 March 2011 – via std.com, pdf scanned copy of the original. Later published as Solomonoff, Ray (1957). "An Inductive Inference Machine". *IRE Convention Record*. Vol. Section on Information Theory, part 2. pp. 56–62.
 - Spadafora, Anthony (21 October 2016). "Stephen Hawking believes AI could be mankind's last accomplishment" (<https://betanews.com/2016/10/21/artificial-intelligence-stephen-hawking/>). BetaNews. Archived (<https://web.archive.org/web/20170828183930/https://betanews.com/2016/10/21/artificial-intelligence-stephen-hawking/>) from the original on 28 August 2017.
 - Tao, Jianhua; Tan, Tieniu (2005). *Affective Computing and Intelligent Interaction*. Affective Computing: A Review. Lecture Notes in Computer Science. Vol. LNCS 3784. Springer. pp. 981–995. doi:[10.1007/11573548](https://doi.org/10.1007/11573548) (<https://doi.org/10.1007%2F11573548>). ISBN 978-3-540-29621-8.
 - Tecuci, Gheorghe (March–April 2012). "Artificial Intelligence". *Wiley Interdisciplinary Reviews: Computational Statistics*. 4 (2): 168–180. doi:[10.1002/wics.200](https://doi.org/10.1002/wics.200) (<https://doi.org/10.1002%2Fwics.200>). S2CID 196141190 (<https://api.semanticscholar.org/CorpusID:196141190>).
 - Thro, Ellen (1993). *Robotics: The Marriage of Computers and Machines* (https://archive.org/details/isbn_9780816026289). New York: Facts on File. ISBN 978-0-8160-2628-9. Archived (https://web.archive.org/web/20200726131505/https://archive.org/details/isbn_9780816026289) from the original on 26 July 2020. Retrieved 22 August 2020.
 - Turing, Alan (October 1950), "Computing Machinery and Intelligence", *Mind*, LIX (236): 433–460, doi:[10.1093/mind/LIX.236.433](https://doi.org/10.1093/mind/LIX.236.433) (<https://doi.org/10.1093%2Fmind%2FLIX.236.433>), ISSN 0026-4423 (<https://www.worldcat.org/issn/0026-4423>).
 - Vinge, Vernor (1993). "The Coming Technological Singularity: How to Survive in the Post-Human Era" (<https://web.archive.org/web/20070101133646/http://www-rohan.sdsu.edu/faculty/vinge/misc/singularity.html>). *Vision 21: Interdisciplinary Science and Engineering in the Era of Cyberspace*: 11. Bibcode:1993vise.nasa...11V (<https://ui.adsabs.harvard.edu/abs/1993vise.nas...11V>). Archived from the original (<http://www-rohan.sdsu.edu/faculty/vinge/misc/singularity.html>) on 1 January 2007. Retrieved 14 November 2011.
 - Wason, P. C.; Shapiro, D. (1966). "Reasoning" (<https://archive.org/details/newhorizonsinpsy0000foss>). In Foss, B. M. (ed.). *New horizons in psychology*. Harmondsworth: Penguin. Archived (<https://web.archive.org/web/20200726131518/https://archive.org/details/newhorizonsinpsy0000foss>) from the original on 26 July 2020. Retrieved 18 November 2019.
 - Weng, J.; McClelland; Pentland, A.; Sporns, O.; Stockman, I.; Sur, M.; Thelen, E. (2001). "Autonomous mental development by robots and animals" (<http://www.cse.msu.edu/dl/SciencePaper.pdf>) (PDF). *Science*. 291 (5504): 599–600. doi:[10.1126/science.291.5504.599](https://doi.org/10.1126/science.291.5504.599) (<https://doi.org/10.1126%2Fscience.291.5504.599>). PMID 11229402 (<https://pubmed.ncbi.nlm.nih.gov/11229402>). S2CID 54131797 (<https://api.semanticscholar.org/CorpusID:54131797>). Archived (<https://web.archive.org/web/20130904235242/http://www.cse.msu.edu/dl/SciencePaper.pdf>) (PDF) from the original on 4 September 2013. Retrieved 4 June 2013 – via msu.edu.

Further reading

- Autor, David H., "Why Are There Still So Many Jobs? The History and Future of Workplace Automation" (2015) 29(3) *Journal of Economic Perspectives* 3.
- Boden, Margaret, *Mind As Machine*, Oxford University Press, 2006.

- Cukier, Kenneth, "Ready for Robo' 98, no. 4 (July/August 2019), pp. 1–10. the Future of AI", *Foreign Affairs*, vol. 98, no. 4 (July/August 2019), pp. 1–10. historian of computing, writes (in what might be called "Dyson's Law") that "Any system simple enough to be understandable will not be complicated enough to behave intelligently, while any system complicated enough to behave intelligently will be too complicated to understand." (p. 197.) Computer scientist Alex Pentland writes: "Current AI machine-learning algorithms are, at their core, dead simple stupid. They work, but they work by brute force." (p. 198.)
- Domingos, Pedro, "Our Digital Doubles: AI will serve our species, not control it", *Scientific American*, vol. 319, no. 3 (September 2018), pp. 88–93.
- Gopnik, Alison, "Making AI More Human: Artificial intelligence has staged a revival by starting to incorporate what we know about how children learn", *Scientific American*, vol. 316, no. 6 (June 2017), pp. 60–65.
- Halpern, Sue, "The Human Costs of AI" (review of Kate Crawford, *Atlas of AI: Power, Politics, and the Planetary Costs of Artificial Intelligence*, Yale University Press, 2021, 327 pp.; Simon Chesterman, *We, the Robots?: Regulating Artificial Intelligence and the Limits of the Law*, Cambridge University Press, 2021, 289 pp.; Keven Roose, *Futureproof: 9 Rules for Humans in the Age of Automation*, Random House, 217 pp.; Erik J. Larson, *The Myth of Artificial Intelligence: Why Computers Can't Think the Way We Do*, Belknap Press / Harvard University Press, 312 pp.), *The New York Review of Books*, vol. LXVIII, no. 16 (21 October 2021), pp. 29–31. "AI training models can replicate entrenched social and cultural biases. [...] Machines only know what they know from the data they have been given. [p. 30.] [A]rtificial general intelligence—machine-based intelligence that matches our own—is beyond the capacity of algorithmic machine learning... 'Your brain is one piece in a broader system which includes your body, your environment, other humans, and culture as a whole.' [E]ven machines that master the tasks they are trained to perform can't jump domains. AIVA, for example, can't drive a car even though it can write music (and wouldn't even be able to do that without Bach and Beethoven [and other composers on which AIVA is trained])." (p. 31.)
- Johnston, John (2008) *The Allure of Machinic Life: Cybernetics, Artificial Life, and the New AI*, MIT Press.
- Koch, Christof, "Proust among the Machines", *Scientific American*, vol. 321, no. 6 (December 2019), pp. 46–49. Christof Koch doubts the possibility of "intelligent" machines attaining consciousness, because "[e]ven the most sophisticated brain simulations are unlikely to produce conscious feelings." (p. 48.) According to Koch, "Whether machines can become sentient [is important] for ethical reasons. If computers experience life through their own senses, they cease to be purely a means to an end determined by their usefulness to... humans. Per GNW [the Global Neuronal Workspace theory], they turn from mere objects into subjects... with a point of view.... Once computers' cognitive abilities rival those of humanity, their impulse to push for legal and political rights will become irresistible—the right not to be deleted, not to have their memories wiped clean, not to suffer pain and degradation. The alternative, embodied by IIT [Integrated Information Theory], is that computers will remain only supersophisticated machinery, ghostlike empty shells, devoid of what we value most: the feeling of life itself." (p. 49.)
- Marcus, Gary, "Am I Human?: Researchers need new ways to distinguish artificial intelligence from the natural kind", *Scientific American*, vol. 316, no. 3 (March 2017), pp. 58–63. A stumbling block to AI has been an incapacity for reliable disambiguation. An example is the "pronoun disambiguation problem": a machine has no way of determining to whom or what a pronoun in a sentence refers. (p. 61.)
- Gary Marcus, "Artificial Confidence: Even the newest,uzziest systems of artificial general intelligence are stymied by the same old problems", *Scientific American*, vol. 327, no. 4 (October 2022), pp. 42–45.
- E McGaughey, 'Will Robots Automate Your Job Away? Full Employment, Basic Income, and Economic Democracy' (2018) SSRN, part 2(3) (https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3044448) Archived (https://web.archive.org/web/20180524201340/https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3044448) 24 May 2018 at the Wayback Machine.

- George Musser, "Artificial Intelligence: What it is, what it's not, and how it will learn creativity and common sense, among other human qualities", *Scientific American*, vol. 320, no. 5 (May 2019), pp. 58–63.
- Myers, Courtney Boyd ed. (2009). "The AI Report" (https://www.forbes.com/2009/06/22/singularity-robots-computers-opinions-contributors-artificial-intelligence-09_land.html) Archived (http://web.archive.org/web/20170729114303/https://www.forbes.com/2009/06/22/singularity-robots-computers-opinions-contributors-artificial-intelligence-09_land.html) 29 July 2017 at the Wayback Machine. *Forbes* June 2009
- Raphael, Bertram (1976). *The Thinking Computer* (<https://archive.org/details/thinkingcomputer00raph>). W.H. Freeman and Co. ISBN 978-0716707233. Archived (<https://web.archive.org/web/20200726215746/https://archive.org/details/thinkingcomputer00raph>) from the original on 26 July 2020. Retrieved 22 August 2020.
- Scharre, Paul, "Killer Apps: The Real Dangers of an AI Arms Race", *Foreign Affairs*, vol. 98, no. 3 (May/June 2019), pp. 135–44. "Today's AI technologies are powerful but unreliable. Rules-based systems cannot deal with circumstances their programmers did not anticipate. Learning systems are limited by the data on which they were trained. AI failures have already led to tragedy. Advanced autopilot features in cars, although they perform well in some circumstances, have driven cars without warning into trucks, concrete barriers, and parked cars. In the wrong situation, AI systems go from supersmart to superdumb in an instant. When an enemy is trying to manipulate and hack an AI system, the risks are even greater." (p. 140.)
- Serenko, Alexander (2010). "The development of an AI journal ranking based on the revealed preference approach" (http://www.aserenko.com/papers/JOI_Serenko_AI_Journal_Ranking_Published.pdf) (PDF). *Journal of Informetrics*. 4 (4): 447–59. doi:10.1016/j.joi.2010.04.001 (<https://doi.org/10.1016%2Fj.joi.2010.04.001>). Archived (https://web.archive.org/web/20131004215236/http://www.aserenko.com/papers/JOI_Serenko_AI_Journal_Ranking_Published.pdf) (PDF) from the original on 4 October 2013. Retrieved 24 August 2013.
- Serenko, Alexander; Michael Dohan (2011). "Comparing the expert survey and citation impact journal ranking methods: Example from the field of Artificial Intelligence" (http://www.aserenko.com/papers/JOI_AI_Journal_Ranking_Serenko.pdf) (PDF). *Journal of Informetrics*. 5 (4): 629–49. doi:10.1016/j.joi.2011.06.002 (<https://doi.org/10.1016%2Fj.joi.2011.06.002>). Archived (https://web.archive.org/web/20131004212839/http://www.aserenko.com/papers/JOI_AI_Journal_Ranking_Serenko.pdf) (PDF) from the original on 4 October 2013. Retrieved 12 September 2013.
- Tom Simonite (29 December 2014). "2014 in Computing: Breakthroughs in Artificial Intelligence" (<http://www.technologyreview.com/news/533686/2014-in-computing-breakthroughs-in-artificial-intelligence/>). *MIT Technology Review*.
- Sun, R. & Bookman, L. (eds.), *Computational Architectures: Integrating Neural and Symbolic Processes*. Kluwer Academic Publishers, Needham, MA. 1994.
- Taylor, Paul, "Insanely Complicated, Hopelessly Inadequate" (review of Brian Cantwell Smith, *The Promise of Artificial Intelligence: Reckoning and Judgment*, MIT, 2019, ISBN 978-0262043045, 157 pp.; Gary Marcus and Ernest Davis, *Rebooting AI: Building Artificial Intelligence We Can Trust*, Ballantine, 2019, ISBN 978-1524748258, 304 pp.; Judea Pearl and Dana Mackenzie, *The Book of Why: The New Science of Cause and Effect*, Penguin, 2019, ISBN 978-0141982410, 418 pp.), *London Review of Books*, vol. 43, no. 2 (21 January 2021), pp. 37–39. Paul Taylor writes (p. 39): "Perhaps there is a limit to what a computer can do without knowing that it is manipulating imperfect representations of an external reality."

- Tooze, Adam, "Democracy and Its 'York Review of Books', vol. LXVI, no. 10 (6 June 2019), pp. 52–53, no clear answer for the mindless operation of bureaucratic and technological power. we may indeed be witnessing its extension in the form of artificial intelligence and robotics. Likewise, after decades of dire warning, the environmental problem remains fundamentally unaddressed.... Bureaucratic overreach and environmental catastrophe are precisely the kinds of slow-moving existential challenges that democracies deal with very badly.... Finally, there is the threat du jour: corporations and the technologies they promote." (pp. 56–57.)

External links

- "Artificial Intelligence" (<http://www.iep.utm.edu/art-inte>). *Internet Encyclopedia of Philosophy*.
 - Thomason, Richmond. "Logic and Artificial Intelligence" (<https://plato.stanford.edu/entries/logic-ai/>). In Zalta, Edward N. (ed.). *Stanford Encyclopedia of Philosophy*.
 - Artificial Intelligence (<https://www.bbc.co.uk/programmes/p003k9fc>). BBC Radio 4 discussion with John Agar, Alison Adam & Igor Aleksander (*In Our Time*, 8 December 2005).
-

Retrieved from "https://en.wikipedia.org/w/index.php?title=Artificial_intelligence&oldid=1112753222"

This page was last edited on 27 September 2022, at 23:05 (UTC).

Text is available under the Creative Commons Attribution-ShareAlike License 3.0; additional terms may apply. By using this site, you agree to the Terms of Use and Privacy Policy. Wikipedia® is a registered trademark of the Wikimedia Foundation, Inc., a non-profit organization.